

Mudcard

- **At the very beginning of the lecture when you were speaking about continuous column visualization, and changing the histogram to show more information, I did not understand how to two below lines of codes changed the histograms**
 - Just try it on the notebook to figure it out and read the histogram manual. :)
- **Are we supposed to use the defaulted figsize for all the graphs?**
 - usually the answer to such questions is no.
 - the default figsize might work most of the time but sometimes you'll have to change it.
 - I do not use the default figsize in the lecture notes.
 - I make the figures smaller than default because I need to increase the font size to make it readable in class
 - you need to experiment to figure out what works
 - e.g., you might need different figure sizes / font sizes, etc. for the report and the presentation
- **when to use logs**
 - if the quantity you plot varies over multiple orders of magnitudes
- **The muddiest part of the lecture was bring able to decide which type of data visualization to use based on whether the variables are continuous, ordinal, or categorical. I feel confident in classifying variables as such, but once I do that, I don't necessarily know which chart to use.**
 - rewatch the video recording and makes notes on this
 - I described 1 to 3 figure types based on column properties
- **How do we make the x-axis log?**
 - `plt.semilogx()`
 - or search the matplotlib manual and/or stackoverflow and/or chatGPT or your favorite generative AI tool
- **Interpretation of Violin Plots and Scatter Matrices**
 - read more about these figure types in the resources I linked at the end of the lecture
- **Heat maps, bins, and changing scales to log were difficult to understand**
 - experiment with the code in the notebook

Split iid data

By the end of this lecture, you will be able to

- apply basic split to iid datasets
- apply k-fold split to iid datasets
- apply stratified splits to imbalanced data

The supervised ML pipeline

The goal: Use the training data (X and y) to develop a **model** which can **accurately** predict the target variable (y_{new}) for previously unseen data (X_{new}).

1. Exploratory Data Analysis (EDA): you need to understand your data and verify that it doesn't contain errors

- do as much EDA as you can!

2. Split the data into different sets: most often the sets are train, validation, and test (or holdout)

- practitioners often make errors in this step!
- you can split the data randomly, based on groups, based on time, or any other non-standard way if necessary to answer your ML question

3. Preprocess the data: ML models only work if X and Y are numbers! Some ML models additionally require each feature to have 0 mean and 1 standard deviation (standardized features)

- often the original features you get contain strings (for example a gender feature would contain 'male', 'female', 'non-binary', 'unknown') which needs to be transformed into numbers
- often the features are not standardized (e.g., age is between 0 and 100) but it needs to be standardized

4. Choose an evaluation metric: depends on the priorities of the stakeholders

- often requires quite a bit of thinking and ethical considerations

5. Choose one or more ML techniques: it is highly recommended that you try multiple models

- start with simple models like linear or logistic regression
- try also more complex models like nearest neighbors, support vector machines, random forest, etc.

6. Tune the hyperparameters of your ML models (aka cross-validation)

- ML techniques have hyperparameters that you need to optimize to achieve best performance
- for each ML model, decide which parameters to tune and what values to try
- loop through each parameter combination
 - train one model for each parameter combination
 - evaluate how well the model performs on the validation set
- take the parameter combo that gives the best validation score
- evaluate that model on the test set to report how well the model is expected to perform on previously unseen data

7. Interpret your model: black boxes are often not useful

- check if your model uses features that make sense (excellent tool for debugging)
- often model predictions are not enough, you need to be able to explain how the model arrived to a particular prediction (e.g., in health care)

Why do we split the data?

- we want to find the best hyper-parameters of our ML algorithms
 - fit models to training data
 - evaluate each model on validation set
 - we find hyper-parameter values that optimize the validation score
- we want to know how the model will perform on previously unseen data
 - apply our final model on the test set

We need to split the data into three parts!

Recap from the second lecture

- **the learner's input**
 - Domain set \mathcal{X} - a set of objects we wish to label.
 - Label set \mathcal{Y} - a set of possible labels.
 - Training data $S = ((x_1, y_1), \dots, (x_m, y_m))$ - a finite sequence of pairs from \mathcal{X}, \mathcal{Y} . This is what the learner has access to.
 - $X = (x_1, \dots, x_m)$ is the feature matrix which is usually a 2D matrix, and $Y = (y_1, \dots, y_m)$ is the target variable which is a vector.
- let's denote the probability distribution over \mathcal{X} by D .
- let's assume there is some correct labeling function $f : \mathcal{X} \rightarrow \mathcal{Y}$.
- **a training example is then generated by sampling x_i from D , and the label y_i is generated using f .**

I.I.D. assumption

- **the i.i.d. assumption:** the examples in the training set are independently and identically distributed according to D
 - every x_i is freshly sampled from D and then labelled by f
 - that is, x_i and y_i are picked independently of the other instances
 - S is a window through which the learner gets partial info about D and the labeling function f
 - the larger the sample gets, the more likely it is to reflect more accurately D and f
- examples of not iid data:
 - data generated by time-dependent processes
 - data has group structure (samples collected from e.g., different subjects, experiments, measurement devices)

Split iid data

By the end of this lecture, you will be able to

- apply basic split to iid datasets
- apply k-fold split to iid datasets
- apply stratified splits to imbalanced data

Splitting strategies for iid data: basic approach

- 60% train, 20% validation, 20% test for small datasets
- 98% train, 1% validation, 1% test for large datasets
 - if you have 1 million points, you still have 10000 points in validation and test which is plenty to assess model performance

Let's work with the adult data!

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

df = pd.read_csv('data/adult_test.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k or less than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features
print(y)
print(X.head())
```

0	<=50K.
1	<=50K.
2	>50K.
3	>50K.
4	<=50K.
	...
16276	<=50K.
16277	<=50K.
16278	<=50K.
16279	<=50K.
16280	>50K.

Name: gross-income, Length: 16281, dtype: object

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	25	Private	226802	11th	7	Never-married	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	160323	Some-college	10	Married-civ-spouse	
4	18	?	103497	Some-college	10	Never-married	

	occupation	relationship	race	sex	capital-gain	\
0	Machine-op-inspct	Own-child	Black	Male	0	
1	Farming-fishing	Husband	White	Male	0	
2	Protective-serv	Husband	White	Male	0	
3	Machine-op-inspct	Husband	Black	Male	7688	
4	?	Own-child	White	Female	0	

	capital-loss	hours-per-week	native-country
0	0	40	United-States
1	0	50	United-States
2	0	40	United-States
3	0	40	United-States
4	0	30	United-States

```
In [2]: help(train_test_split)
```

Help on function train_test_split in module sklearn.model_selection._split:

train_test_split(*arrays, test_size=None, train_size=None, random_state=None, shuffle=True, stratify=None)
Split arrays or matrices into random train and test subsets.

Quick utility that wraps input validation, `next(ShuffleSplit().split(X, y))`, and application to input data into a single call for splitting (and optionally subsampling) data into a one-liner.

Read more in the :ref:`User Guide <cross_validation>`.

Parameters

`*arrays` : sequence of indexables with same length / shape[0]

Allowed inputs are lists, numpy arrays, scipy-sparse matrices or pandas dataframes.

`test_size` : float or int, default=None

If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the test split. If int, represents the absolute number of test samples. If None, the value is set to the complement of the train size. If `train_size` is also None, it will be set to 0.25.

`train_size` : float or int, default=None

If float, should be between 0.0 and 1.0 and represent the proportion of the dataset to include in the train split. If int, represents the absolute number of train samples. If None, the value is automatically set to the complement of the test size.

`random_state` : int, RandomState instance or None, default=None

Controls the shuffling applied to the data before applying the split. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random_state>`.

`shuffle` : bool, default=True

Whether or not to shuffle the data before splitting. If `shuffle=False` then `stratify` must be None.

`stratify` : array-like, default=None

If not None, data is split in a stratified fashion, using this as the class labels.
Read more in the :ref:`User Guide <stratification>`.

Returns

`splitting` : list, length=2 * len(arrays)
List containing train-test split of inputs.

.. versionadded:: 0.16

If the input is sparse, the output will be a `scipy.sparse.csr_matrix`. Else, output type is the same as the input type.

Examples

```
>>> import numpy as np
>>> from sklearn.model_selection import train_test_split
>>> X, y = np.arange(10).reshape((5, 2)), range(5)
>>> X
array([[0, 1],
       [2, 3],
       [4, 5],
       [6, 7],
       [8, 9]])
>>> list(y)
[0, 1, 2, 3, 4]

>>> X_train, X_test, y_train, y_test = train_test_split(
...     X, y, test_size=0.33, random_state=42)
...
>>> X_train
array([[4, 5],
       [0, 1],
       [6, 7]])
>>> y_train
[2, 0, 3]
>>> X_test
```

```

array([[2, 3],
       [8, 9]])
>>> y_test
[1, 4]

>>> train_test_split(y, shuffle=False)
[[0, 1, 2], [3, 4]]

```

```

In [3]: random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,\
                                                    train_size = 0.6,random_state = random_state)
print('training set:',X_train.shape, y_train.shape) # 60% of points are in train
print(X_other.shape, y_other.shape) # 40% of points are in other

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,\
                                                train_size = 0.5,random_state = random_state)
print('validation set:',X_val.shape, y_val.shape) # 20% of points are in validation
print('test set:',X_test.shape, y_test.shape) # 20% of points are in test

print(X_train.head())

```

```

training set: (9768, 14) (9768,)
(6513, 14) (6513,)
validation set: (3256, 14) (3256,)
test set: (3257, 14) (3257,)

```

	age	workclass	fnlwgt	education	education-num	\
4050	22	Private	335950	HS-grad	9	
11446	29	Private	78261	HS-grad	9	
12427	74	Self-emp-not-inc	160009	Assoc-acdm	12	
5702	39	Self-emp-inc	31709	Some-college	10	
13058	50	Private	144084	HS-grad	9	

	marital-status	occupation	relationship	race	sex	\
4050	Never-married	Other-service	Not-in-family	Black	Male	
11446	Separated	Protective-serv	Not-in-family	White	Male	
12427	Married-civ-spouse	Exec-managerial	Husband	White	Male	
5702	Married-civ-spouse	Adm-clerical	Wife	White	Female	
13058	Separated	Sales	Unmarried	White	Female	

	capital-gain	capital-loss	hours-per-week	native-country
4050	0	0	70	United-States
11446	0	0	55	United-States
12427	0	0	30	United-States
5702	0	0	20	United-States
13058	0	0	55	United-States

Randomness due to splitting

- the model performance, validation and test scores will change depending on which points are in train, val, test
 - inherent randomness or uncertainty of the ML pipeline
- change the random state a couple of times and repeat the whole ML pipeline to assess how much the random splitting affects your test score
 - you would expect a similar uncertainty when the model is deployed

Quiz 1

What's the second `train_test_split` line if you want to end up with 60-20-20 in train-val-test? Print out the sizes of `X_train`, `X_val`, `X_test` to verify!

```

In [4]: X_other, X_test, y_other, y_test = train_test_split(X,y,\
                                                         train_size = 0.8,random_state=random_state)
# add your line below and chose the correct solution from canvas

```

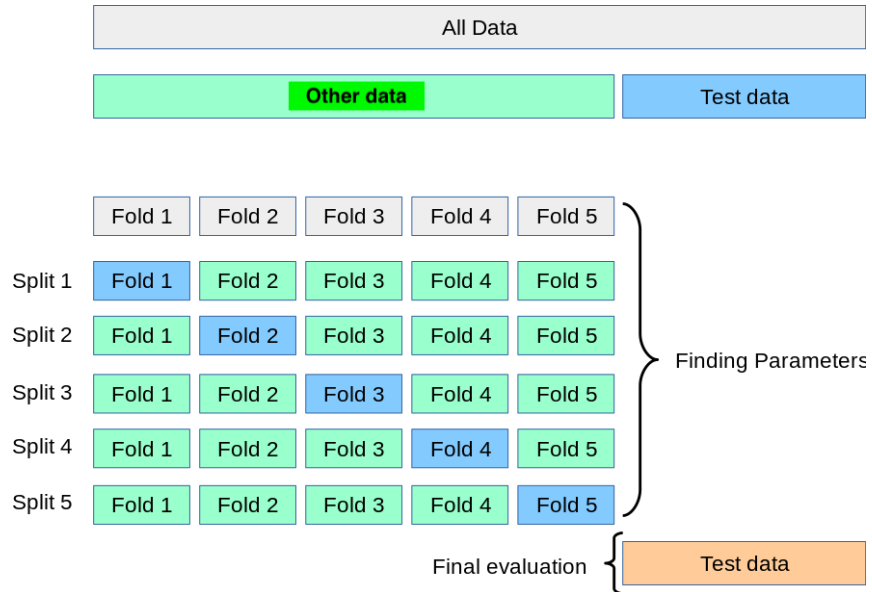
Split iid data

By the end of this lecture, you will be able to

- apply basic split to iid datasets

- apply k-fold split to iid datasets
- apply stratified splits to imbalanced data

Other splitting strategy for iid data: k-fold splitting



```
In [5]: from sklearn.model_selection import KFold
        help(KFold)
```

Help on class KFold in module sklearn.model_selection._split:

```
class KFold(_BaseKFold)
|   KFold(n_splits=5, *, shuffle=False, random_state=None)
|
|   K-Folds cross-validator
|
|   Provides train/test indices to split data in train/test sets. Split
|   dataset into k consecutive folds (without shuffling by default).
|
|   Each fold is then used once as a validation while the k - 1 remaining
|   folds form the training set.
|
|   Read more in the :ref:`User Guide <k_fold>`.
|
|   Parameters
|   -----
|   n_splits : int, default=5
|       Number of folds. Must be at least 2.
|
|       .. versionchanged:: 0.22
|           ``n_splits`` default value changed from 3 to 5.
|
|   shuffle : bool, default=False
|       Whether to shuffle the data before splitting into batches.
|       Note that the samples within each split will not be shuffled.
|
|   random_state : int, RandomState instance or None, default=None
|       When `shuffle` is True, `random_state` affects the ordering of the
|       indices, which controls the randomness of each fold. Otherwise, this
|       parameter has no effect.
|       Pass an int for reproducible output across multiple function calls.
|       See :term:`Glossary <random_state>`.
|
|   Examples
|   -----
|   >>> import numpy as np
|   >>> from sklearn.model_selection import KFold
|   >>> X = np.array([[1, 2], [3, 4], [1, 2], [3, 4]])
|   >>> y = np.array([1, 2, 3, 4])
|   >>> kf = KFold(n_splits=2)
|   >>> kf.get_n_splits(X)
|   2
|   >>> print(kf)
|   KFold(n_splits=2, random_state=None, shuffle=False)
|   >>> for i, (train_index, test_index) in enumerate(kf.split(X)):
|   ...     print(f"Fold {i}:")
|   ...     print(f"  Train: index={train_index}")
|   ...     print(f"  Test:  index={test_index}")
|   Fold 0:
|     Train: index=[2 3]
|     Test:  index=[0 1]
|   Fold 1:
|     Train: index=[0 1]
|     Test:  index=[2 3]
|
|   Notes
|   -----
|   The first ``n_samples % n_splits`` folds have size
|   ``n_samples // n_splits + 1``, other folds have size
|   ``n_samples // n_splits``, where ``n_samples`` is the number of samples.
|
|   Randomized CV splitters may return different results for each call of
|   split. You can make the results identical by setting `random_state`
|   to an integer.
|
|   See Also
|   -----
|   StratifiedKFold : Takes class information into account to avoid building
|       folds with imbalanced class distributions (for binary or multiclass
|       classification tasks).
|
|   GroupKFold : K-fold iterator variant with non-overlapping groups.
|
|   RepeatedKFold : Repeats K-Fold n times.
|
|   Method resolution order:
|       KFold
|       _BaseKFold
```

```

BaseCrossValidator
sklearn.utils._metadata_requests._MetadataRequester
builtins.object

Methods defined here:

__init__(self, n_splits=5, *, shuffle=False, random_state=None)
    Initialize self. See help(type(self)) for accurate signature.

-----

Data and other attributes defined here:

__abstractmethods__ = frozenset()

-----

Methods inherited from _BaseKFold:

get_n_splits(self, X=None, y=None, groups=None)
    Returns the number of splitting iterations in the cross-validator

    Parameters
    -----
    X : object
        Always ignored, exists for compatibility.

    y : object
        Always ignored, exists for compatibility.

    groups : object
        Always ignored, exists for compatibility.

    Returns
    -----
    n_splits : int
        Returns the number of splitting iterations in the cross-validator.

split(self, X, y=None, groups=None)
    Generate indices to split data into training and test set.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Training data, where `n_samples` is the number of samples
        and `n_features` is the number of features.

    y : array-like of shape (n_samples,), default=None
        The target variable for supervised learning problems.

    groups : array-like of shape (n_samples,), default=None
        Group labels for the samples used while splitting the dataset into
        train/test set.

    Yields
    -----
    train : ndarray
        The training set indices for that split.

    test : ndarray
        The testing set indices for that split.

-----

Methods inherited from BaseCrossValidator:

__repr__(self)
    Return repr(self).

-----

Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

get_metadata_routing(self)
    Get metadata routing of this object.

    Please check :ref:`User Guide <metadata_routing>` on how the routing
    mechanism works.

    Returns
    -----
    routing : MetadataRequest
        A :class:`~utils.metadata_routing.MetadataRequest` encapsulating

```


routing information.

Class methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

`__init_subclass__`(*kwargs) from abc.ABCMeta
Set the `__set_{method}_request__` methods.

This uses PEP-487 [1] to set the `__set_{method}_request__` methods. It looks for the information available in the set default values which are set using `__metadata_request__` class attributes, or inferred from method signatures.

The `__metadata_request__` class attributes are used when a method does not explicitly accept a metadata through its arguments or if the developer would like to specify a request value for those metadata which are different from the default `None`.

References

.. [1] <https://www.python.org/dev/peps/pep-0487>

Data descriptors inherited from sklearn.utils._metadata_requests._MetadataRequester:

`__dict__`
dictionary for instance variables (if defined)

`__weakref__`
list of weak references to the object (if defined)

```
In [6]: random_state = 42

# first split to separate out the test set
X_other, X_test, y_other, y_test = train_test_split(X, y, test_size = 0.2, random_state=random_state)
print(X_other.shape, y_other.shape)
print('test set:', X_test.shape, y_test.shape)

# do KFold split on other
kf = KFold(n_splits=5, shuffle=True, random_state=random_state)
for train_index, val_index in kf.split(X_other, y_other):
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print('  training set:', X_train.shape, y_train.shape)
    print('  validation set:', X_val.shape, y_val.shape)
    # the validation set contains different points in each iteration
    print(X_val[['age', 'workclass', 'education']].head())
```

```

(13024, 14) (13024,)
test set: (3257, 14) (3257,)
  training set: (10419, 14) (10419,)
  validation set: (2605, 14) (2605,)
    age      workclass      education
9850    59      Private    Some-college
103     58    Self-emp-not-inc    9th
1383    45      Private      HS-grad
11034   49    Self-emp-not-inc    Bachelors
14876   59    Self-emp-not-inc    Bachelors
  training set: (10419, 14) (10419,)
  validation set: (2605, 14) (2605,)
    age      workclass      education
13384   60    Federal-gov    Bachelors
8471    20      Private      HS-grad
13406   21      ?    Some-college
13394   35      Private      HS-grad
15123   38      Private    Some-college
  training set: (10419, 14) (10419,)
  validation set: (2605, 14) (2605,)
    age      workclass      education
647     60      ?    Bachelors
9314    26      Private    Some-college
14499   52      Private      HS-grad
7332    53    Federal-gov    Assoc-acdm
12523   21      Private      10th
  training set: (10419, 14) (10419,)
  validation set: (2605, 14) (2605,)
    age      workclass      education
5294    53      Private      HS-grad
3481    41      Private      HS-grad
7671    49      Private    Some-college
11055   39      Private    Bachelors
12751   18      ?      12th
  training set: (10420, 14) (10420,)
  validation set: (2604, 14) (2604,)
    age      workclass      education
4265    23      ?      10th
5290    23      Private      HS-grad
1157    56    Self-emp-inc    Prof-school
12344   18      Private      11th
13683   55      Private      HS-grad


```

How many splits should I create?

- tough question, 3-5 is most common
- if you do n splits, n models will be trained, so the larger the n , the most computationally intensive it will be to train the models
- KFold is usually better suited to small datasets
- KFold is good to estimate uncertainty due to random splitting of train and val, but it is not perfect
 - the test set remains the same

Why shuffling iid data is important?

- by default, data is not shuffled by Kfold which can introduce errors!



Quiz 2

Given the labels below, what are the balances of each class?

y = [0,0,0,2,2,0,0,2,0,1]

Split iid data

By the end of this lecture, you will be able to

- apply basic split to iid datasets
- apply k-fold split to iid datasets
- **apply stratified splits to imbalanced data**

Imbalanced data

- imbalanced data: only a small fraction of the points are in one of the classes, usually ~5% or less but there is no hard limit here
- examples:
 - people visit a bank's website. do they sign up for a new credit card?
 - most customers just browse and leave the page
 - usually 1% or less of the customers get a credit card (class 1), the rest leaves the page without signing up (class 0).
 - fraud detection
 - only a tiny fraction of credit card payments are fraudulent
 - rare disease diagnosis
- the issue with imbalanced data:
 - if you apply train_test_split or KFold, you might not have class 1 points in one of your sets by chance
 - this is what we need to fix

Solution: stratified splits

```
In [7]: df = pd.read_csv('data/imbalanced_data.csv')

X = df[['feature1', 'feature2']]
y = df['y']

print(y.value_counts())
```

```
y
0    990
1     10
Name: count, dtype: int64
```

```
In [8]: # 4 and 10

random_state = 4

X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.6,random_state=random_state)
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.5,random_state=random_state)

print('**balance without stratification:**')
# a variation on the order of 1% which would be too much for imbalanced data!
print(np.unique(y_train,return_counts=True))
print(np.unique(y_val,return_counts=True))
print(np.unique(y_test,return_counts=True))


X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.6,stratify=y,random_state=random_state)
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.5,stratify=y_other,random_state=random_state)
print('**balance with stratification:**')
# very little variation (in the 4th decimal point only) which is important if the problem is imbalanced
print(np.unique(y_train,return_counts=True))
print(np.unique(y_val,return_counts=True))
print(np.unique(y_test,return_counts=True))
```

```

**balance without stratification:**
(array([0, 1]), array([591, 9]))
(array([0]), array([200]))
(array([0, 1]), array([199, 1]))
**balance with stratification:**
(array([0, 1]), array([594, 6]))
(array([0, 1]), array([198, 2]))
(array([0, 1]), array([198, 2]))

```

Stratified folds



```

In [9]: from sklearn.model_selection import StratifiedKFold
        help(StratifiedKFold)

```

Help on class StratifiedKFold in module sklearn.model_selection._split:

```
class StratifiedKFold(_BaseKFold)
|   StratifiedKFold(n_splits=5, *, shuffle=False, random_state=None)
|
|   Stratified K-Folds cross-validator.
|
|   Provides train/test indices to split data in train/test sets.
|
|   This cross-validation object is a variation of KFold that returns
|   stratified folds. The folds are made by preserving the percentage of
|   samples for each class.
|
|   Read more in the :ref:`User Guide <stratified_k_fold>`.
|
|   Parameters
|   -----
|   n_splits : int, default=5
|       Number of folds. Must be at least 2.
|
|       .. versionchanged:: 0.22
|           ``n_splits`` default value changed from 3 to 5.
|
|   shuffle : bool, default=False
|       Whether to shuffle each class's samples before splitting into batches.
|       Note that the samples within each split will not be shuffled.
|
|   random_state : int, RandomState instance or None, default=None
|       When `shuffle` is True, `random_state` affects the ordering of the
|       indices, which controls the randomness of each fold for each class.
|       Otherwise, leave `random_state` as `None`.
|       Pass an int for reproducible output across multiple function calls.
|       See :term:`Glossary <random_state>`.
|
|   Examples
|   -----
|   >>> import numpy as np
|   >>> from sklearn.model_selection import StratifiedKFold
|   >>> X = np.array([[1, 2], [3, 4], [1, 2], [3, 4]])
|   >>> y = np.array([0, 0, 1, 1])
|   >>> skf = StratifiedKFold(n_splits=2)
|   >>> skf.get_n_splits(X, y)
|   2
|   >>> print(skf)
|   StratifiedKFold(n_splits=2, random_state=None, shuffle=False)
|   >>> for i, (train_index, test_index) in enumerate(skf.split(X, y)):
|   ...     print(f"Fold {i}:")
|   ...     print(f"  Train: index={train_index}")
|   ...     print(f"  Test:  index={test_index}")
|   Fold 0:
|   Train: index=[1 3]
|   Test:  index=[0 2]
|   Fold 1:
|   Train: index=[0 2]
|   Test:  index=[1 3]
|
|   Notes
|   -----
|   The implementation is designed to:
|
|   * Generate test sets such that all contain the same distribution of
|     classes, or as close as possible.
|   * Be invariant to class label: relabelling ``y = ["Happy", "Sad"]`` to
|     ``y = [1, 0]`` should not change the indices generated.
|   * Preserve order dependencies in the dataset ordering, when
|     ``shuffle=False``: all samples from class k in some test set were
|     contiguous in y, or separated in y by samples from classes other than k.
|   * Generate test sets where the smallest and largest differ by at most one
|     sample.
|
|       .. versionchanged:: 0.22
|           The previous implementation did not follow the last constraint.
|
|   See Also
|   -----
|   RepeatedStratifiedKFold : Repeats Stratified K-Fold n times.
|
|   Method resolution order:
|   StratifiedKFold
```

```

_BaseKFold
BaseCrossValidator
sklearn.utils._metadata_requests._MetadataRequester
builtins.object

Methods defined here:

__init__(self, n_splits=5, *, shuffle=False, random_state=None)
    Initialize self. See help(type(self)) for accurate signature.

split(self, X, y, groups=None)
    Generate indices to split data into training and test set.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Training data, where `n_samples` is the number of samples
        and `n_features` is the number of features.

        Note that providing ``y`` is sufficient to generate the splits and
        hence ``np.zeros(n_samples)`` may be used as a placeholder for
        ``X`` instead of actual training data.

    y : array-like of shape (n_samples,)
        The target variable for supervised learning problems.
        Stratification is done based on the y labels.

    groups : object
        Always ignored, exists for compatibility.

    Yields
    -----
    train : ndarray
        The training set indices for that split.

    test : ndarray
        The testing set indices for that split.

    Notes
    -----
    Randomized CV splitters may return different results for each call of
    split. You can make the results identical by setting `random_state`
    to an integer.

```

```

Data and other attributes defined here:

__abstractmethods__ = frozenset()

```

```

Methods inherited from _BaseKFold:

get_n_splits(self, X=None, y=None, groups=None)
    Returns the number of splitting iterations in the cross-validator

    Parameters
    -----
    X : object
        Always ignored, exists for compatibility.

    y : object
        Always ignored, exists for compatibility.

    groups : object
        Always ignored, exists for compatibility.

    Returns
    -----
    n_splits : int
        Returns the number of splitting iterations in the cross-validator.

```

```

Methods inherited from BaseCrossValidator:

__repr__(self)
    Return repr(self).

```

```

Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

```

```

| get_metadata_routing(self)
|     Get metadata routing of this object.
|
|     Please check :ref:`User Guide <metadata_routing>` on how the routing
|     mechanism works.
|
|     Returns
|     -----
|     routing : MetadataRequest
|         A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
|         routing information.
|
| -----
| Class methods inherited from sklearn.utils._metadata_requests._MetadataRequester:
|
| __init_subclass__(**kwargs) from abc.ABCMeta
|     Set the ``set_{method}_request`` methods.
|
|     This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It
|     looks for the information available in the set default values which are
|     set using ``__metadata_request_*`` class attributes, or inferred
|     from method signatures.
|
|     The ``__metadata_request_*`` class attributes are used when a method
|     does not explicitly accept a metadata through its arguments or if the
|     developer would like to specify a request value for those metadata
|     which are different from the default ``None``.
|
|     References
|     -----
|     .. [1] https://www.python.org/dev/peps/pep-0487
|
| -----
| Data descriptors inherited from sklearn.utils._metadata_requests._MetadataRequester:
|
| __dict__
|     dictionary for instance variables (if defined)
|
| __weakref__
|     list of weak references to the object (if defined)

```

```

In [10]: # what we did before: variance in balance on the order of 1%
random_state = 2

X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.2,random_state=random_state)
print('test balance:',np.unique(y_test,return_counts=True))

# do KFold split on other
kf = KFold(n_splits=4,shuffle=True,random_state=random_state)
for train_index, val_index in kf.split(X_other,y_other):
    print('new fold')
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print(np.unique(y_train,return_counts=True))
    print(np.unique(y_val,return_counts=True))

test balance: (array([0, 1]), array([198,  2]))
new fold
(array([0, 1]), array([596,  4]))
(array([0, 1]), array([196,  4]))
new fold
(array([0, 1]), array([593,  7]))
(array([0, 1]), array([199,  1]))
new fold
(array([0, 1]), array([592,  8]))
(array([0]), array([200]))
new fold
(array([0, 1]), array([595,  5]))
(array([0, 1]), array([197,  3]))

```

```

In [11]: # stratified K Fold: variation in balance is very small (4th decimal point)
random_state = 42

# stratified train-test split
X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.2,stratify=y,random_state=random_state)

```

```

print('test balance:', np.unique(y_test, return_counts=True))

# do StratifiedKFold split on other
kf = StratifiedKFold(n_splits=4, shuffle=True, random_state=random_state)
for train_index, val_index in kf.split(X_other, y_other):
    print('new fold')
    X_train = X_other.iloc[train_index]
    y_train = y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = y_other.iloc[val_index]
    print(np.unique(y_train, return_counts=True))
    print(np.unique(y_val, return_counts=True))

```

```

test balance: (array([0, 1]), array([198,  2]))
new fold
(array([0, 1]), array([594,  6]))
(array([0, 1]), array([198,  2]))
new fold
(array([0, 1]), array([594,  6]))
(array([0, 1]), array([198,  2]))
new fold
(array([0, 1]), array([594,  6]))
(array([0, 1]), array([198,  2]))
new fold
(array([0, 1]), array([594,  6]))
(array([0, 1]), array([198,  2]))

```

Mudcard

In []: