

Mudcard

- **The difference between groupkfold and groupshufflesplit wasn't fully clear**
 - read their manuals and experiment with them using the simple toy datasets in the lecture notes
- **For the seizure example, I would assume that for a single patient a feature like skin temp difference (vs. some non-seizure normal) would be more useful than skin temp (absolute) in predicting type of seizure since different patients probably have different resting temps. Since you were only using absolute skin temp, would that just make that feature less helpful in the ML prediction or unusable?**
 - keep in mind that this is skin temperature so it fluctuates a lot.
 - I did check this and there was no statistically significant difference between skin temps during and before/after a seizure
 - the skin temperature depends on things like how strong ventilation is in the room, is the patient's hand under or above the cover, etc.
- **The muddiest part of the lecture for me was the actual implementation of the splitting algorithms. It seemed that within the splitting function we were creating multiple splitting objects, so I would love to break down what each of those steps accomplished.**
 - yep, go ahead and do it
 - this will be a great learning experience!
- **If you didn't split the time of the seizure for each patient done to increase the dataset size, would you still need to do the kfold?**
 - I assume you mean groupKfold
 - yes because there would still be multiple seizures from the same patient
 - some patients had 5+ seizures during their hospital visit
- **Is the autocorrelation graph used in the future to help split time series data? I want just not understanding how the autocorrelation graph fit into the splitting data conversation.**
 - it will help you to decide how many autoregression features to generate
 - you will see in today's lecture
- **The differences between GroupKFold and GroupShuffleSplit and when to use each of them.**
 - read the manuals
 - as I said in class, I'd use GroupKFold if you have a small number of groups and I'd use GroupShuffleFold if you have a large number of groups

Data preprocessing

By the end of this lecture, you will be able to

- apply one-hot encoding or ordinal encoding to categorical variables
- apply scaling and normalization to continuous variables

The supervised ML pipeline

The goal: Use the training data (X and y) to develop a **model** which can **accurately** predict the target variable (y_new') for previously unseen data (X_new).

1. Exploratory Data Analysis (EDA): you need to understand your data and verify that it doesn't contain errors

- do as much EDA as you can!

2. Split the data into different sets: most often the sets are train, validation, and test (or holdout)

- practitioners often make errors in this step!
- you can split the data randomly, based on groups, based on time, or any other non-standard way if necessary to answer your ML question

3. Preprocess the data: ML models only work if X and Y are numbers! Some ML models additionally require each feature to have 0 mean and 1 standard deviation (standardized features)

- often the original features you get contain strings (for example a gender feature would contain 'male', 'female', 'non-binary', 'unknown') which needs to be transformed into numbers
- often the features are not standardized (e.g., age is between 0 and 100) but it needs to be standardized

4. Choose an evaluation metric: depends on the priorities of the stakeholders

- often requires quite a bit of thinking and ethical considerations

5. Choose one or more ML techniques: it is highly recommended that you try multiple models

- start with simple models like linear or logistic regression
- try also more complex models like nearest neighbors, support vector machines, random forest, etc.

6. Tune the hyperparameters of your ML models (aka cross-validation)

- ML techniques have hyperparameters that you need to optimize to achieve best performance
- for each ML model, decide which parameters to tune and what values to try
- loop through each parameter combination
 - train one model for each parameter combination
 - evaluate how well the model performs on the validation set
- take the parameter combo that gives the best validation score
- evaluate that model on the test set to report how well the model is expected to perform on previously unseen data

7. Interpret your model: black boxes are often not useful

- check if your model uses features that make sense (excellent tool for debugging)
- often model predictions are not enough, you need to be able to explain how the model arrived to a particular prediction (e.g., in health care)

Problem description, why preprocessing is necessary

Data format suitable for ML: 2D numerical values.

X	feature_1	feature_2	...	feature_j	...	feature_m	y
data_point_1	x_11	x_12	...	x_1j	...	x_1m	y_1
data_point_2	x_21	x_22	...	x_2j	...	x_2m	y_2
...
data_point_i	x_i1	x_i2	...	x_ij	...	x_im	y_i
...
data_point_n	x_n1	x_n2	...	x_nj	...	x_nm	y_n

Data almost never comes in a format that's directly usable in ML.

- let's check the adult data

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k or less than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.6,random_state=random_state)

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.5,random_state=random_state)

print('training set')
print(X_train.head()) # lots of strings!
print(y_train.head()) # even our labels are strings and not numbers!
```

```

training set
  age  workclass  fnlwgt  education  education-num  \
25823  31    Private  87418    Assoc-voc           11
10274  41    Private 121718    Some-college      10
27652  61    Private  79827      HS-grad           9
13941  33  State-gov 156015    Bachelors         13
31384  38    Private 167882    Some-college      10

  marital-status  occupation  relationship  race  \
25823  Married-civ-spouse  Exec-managerial    Husband  White
10274  Married-civ-spouse   Craft-repair    Husband  White
27652  Married-civ-spouse  Exec-managerial    Husband  White
13941  Married-civ-spouse  Exec-managerial    Husband  White
31384      Widowed      Other-service  Other-relative  Black

  sex  capital-gain  capital-loss  hours-per-week  native-country
25823  Male           0             0             40    United-States
10274  Male           0             0             40           Italy
27652  Male           0             0             50    United-States
13941  Male           0             0             40    United-States
31384  Female         0             0             45           Haiti
25823  <=50K
10274  <=50K
27652  <=50K
13941  >50K
31384  <=50K
Name: gross-income, dtype: object

```

scikit-learn transformers to the rescue!

Preprocessing is done with various transformers. All transformers have three methods:

- **fit** method: estimates parameters necessary to do the transformation,
- **transform** method: transforms the data based on the estimated parameters,
- **fit_transform** method: both steps are performed at once, this can be faster than doing the steps separately.

Transformers we cover today

- **OneHotEncoder** - converts categorical features into dummy arrays
- **OrdinalEncoder** - converts categorical features into an integer array
- **MinMaxScaler** - scales continuous variables to be between 0 and 1
- **StandardScaler** - standardizes continuous features by removing the mean and scaling to unit variance

By the end of this lecture, you will be able to

- **apply one-hot encoding or ordinal encoding to categorical variables**
- **apply scaling and normalization to continuous variables**

Ordered categorical data: OrdinalEncoder

- use it on categorical features if the categories can be ranked or ordered
 - educational level in the adult dataset
 - reaction to medication is described by words like 'severe', 'no response', 'excellent'
 - any time you know that the categories can be clearly ranked

```

In [2]: from sklearn.preprocessing import OrdinalEncoder
        help(OrdinalEncoder)

```

Help on class OrdinalEncoder in module sklearn.preprocessing._encoders:

```
class OrdinalEncoder(sklearn.base.OneToOneFeatureMixin, _BaseEncoder)
| OrdinalEncoder(*, categories='auto', dtype=<class 'numpy.float64'>, handle_unknown='error', unknown_value=None, encoded_missing_value=nan, min_frequency=None, max_categories=None)
```

Encode categorical features as an integer array.

The input to this transformer should be an array-like of integers or strings, denoting the values taken on by categorical (discrete) features. The features are converted to ordinal integers. This results in a single column of integers (0 to `n_categories - 1`) per feature.

Read more in the :ref:`User Guide <preprocessing_categorical_features>`.

.. versionadded:: 0.20

Parameters

`categories` : 'auto' or a list of array-like, default='auto'
Categories (unique values) per feature:

- 'auto' : Determine categories automatically from the training data.
- list : ``categories[i]`` holds the categories expected in the *i*th column. The passed categories should not mix strings and numeric values, and should be sorted in case of numeric values.

The used categories can be found in the ``categories_`` attribute.

`dtype` : number type, default=np.float64
Desired dtype of output.

`handle_unknown` : {'error', 'use_encoded_value'}, default='error'
When set to 'error' an error will be raised in case an unknown categorical feature is present during transform. When set to 'use_encoded_value', the encoded value of unknown categories will be set to the value given for the parameter 'unknown_value'. In :meth:`inverse_transform`, an unknown category will be denoted as None.

.. versionadded:: 0.24

`unknown_value` : int or np.nan, default=None
When the parameter `handle_unknown` is set to 'use_encoded_value', this parameter is required and will set the encoded value of unknown categories. It has to be distinct from the values used to encode any of the categories in 'fit'. If set to np.nan, the 'dtype' parameter must be a float dtype.

.. versionadded:: 0.24

`encoded_missing_value` : int or np.nan, default=np.nan
Encoded value of missing categories. If set to 'np.nan', then the 'dtype' parameter must be a float dtype.

.. versionadded:: 1.1

`min_frequency` : int or float, default=None
Specifies the minimum frequency below which a category will be considered infrequent.

- If 'int', categories with a smaller cardinality will be considered infrequent.

- If 'float', categories with a smaller cardinality than '`min_frequency * n_samples`' will be considered infrequent.

.. versionadded:: 1.3

Read more in the :ref:`User Guide <encoder_infrequent_categories>`.

`max_categories` : int, default=None
Specifies an upper limit to the number of output categories for each input feature when considering infrequent categories. If there are infrequent categories, 'max_categories' includes the category representing the infrequent categories along with the frequent categories. If 'None', there is no limit to the number of output features.

'max_categories' do **not** take into account missing or unknown categories. Setting 'unknown_value' or 'encoded_missing_value' to an integer will increase the number of unique integer codes by one each.

This can result in up to ``max_categories + 2`` integer codes.

.. versionadded:: 1.3

Read more in the :ref:`User Guide <encoder_infrequent_categories>`.

Attributes

`categories_` : list of arrays

The categories of each feature determined during ```fit``` (in order of the features in `X` and corresponding with the output of ```transform```). This does not include categories that weren't seen during ```fit```.

`n_features_in_` : int

Number of features seen during :term:`fit`.

.. versionadded:: 1.0

`feature_names_in_` : ndarray of shape (``n_features_in_`` ,)

Names of features seen during :term:`fit`. Defined only when ``X`` has feature names that are all strings.

.. versionadded:: 1.0

`infrequent_categories_` : list of ndarray

Defined only if infrequent categories are enabled by setting ``min_frequency`` or ``max_categories`` to a non-default value. ``infrequent_categories_[i]`` are the infrequent categories for feature ``i``. If the feature ``i`` has no infrequent categories ``infrequent_categories_[i]`` is `None`.

.. versionadded:: 1.3

See Also

`OneHotEncoder` : Performs a one-hot encoding of categorical features. This encoding is suitable for low to medium cardinality categorical variables, both in supervised and unsupervised settings.

`TargetEncoder` : Encodes categorical features using supervised signal in a classification or regression pipeline. This encoding is typically suitable for high cardinality categorical variables.

`LabelEncoder` : Encodes target labels with values between 0 and ``n_classes-1``.

Notes

With a high proportion of ``nan`` values, inferring categories becomes slow with Python versions before 3.10. The handling of ``nan`` values was improved from Python 3.10 onwards, (c.f. ``bpo-43475`` <<https://github.com/python/cpython/issues/87641>>`_`).

Examples

Given a dataset with two features, we let the encoder find the unique values per feature and transform the data to an ordinal encoding.

```
>>> from sklearn.preprocessing import OrdinalEncoder
>>> enc = OrdinalEncoder()
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
>>> enc.fit(X)
OrdinalEncoder()
>>> enc.categories_
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> enc.transform([['Female', 3], ['Male', 1]])
array([[0., 2.],
       [1., 0.]])

>>> enc.inverse_transform([[1, 0], [0, 1]])
array(['Male', 1],
      ['Female', 2]], dtype=object)
```

By default, :class:`OrdinalEncoder` is lenient towards missing values by propagating them.

```
>>> import numpy as np
>>> X = [['Male', 1], ['Female', 3], ['Female', np.nan]]
>>> enc.fit_transform(X)
array([[ 1.,  0.],
       [ 0.,  1.],
       [ 0., nan]])
```

You can use the parameter `encoded_missing_value` to encode missing values.

```
>>> enc.set_params(encoded_missing_value=-1).fit_transform(X)
array([[ 1.,  0.],
       [ 0.,  1.],
       [ 0., -1.]])
```

Infrequent categories are enabled by setting `max_categories` or `min_frequency`. In the following example, "a" and "d" are considered infrequent and grouped together into a single category, "b" and "c" are their own categories, unknown values are encoded as 3 and missing values are encoded as 4.

```
>>> X_train = np.array(
...     [[ "a" ] * 5 + [ "b" ] * 20 + [ "c" ] * 10 + [ "d" ] * 3 + [ np.nan ]],
...     dtype=object).T
>>> enc = OrdinalEncoder(
...     handle_unknown="use_encoded_value", unknown_value=3,
...     max_categories=3, encoded_missing_value=4)
>>> _ = enc.fit(X_train)
>>> X_test = np.array([ [ "a" ], [ "b" ], [ "c" ], [ "d" ], [ "e" ], [ np.nan ] ], dtype=object)
>>> enc.transform(X_test)
array([[2.],
       [0.],
       [1.],
       [2.],
       [3.],
       [4.]])
```

Method resolution order:

```
OrdinalEncoder
sklearn.base.OneToOneFeatureMixin
_BaseEncoder
sklearn.base.TransformerMixin
sklearn.utils._set_output._SetOutputMixin
sklearn.base.BaseEstimator
sklearn.utils._metadata_requests._MetadataRequester
builtins.object
```

Methods defined here:

```
__init__(self, *, categories='auto', dtype=<class 'numpy.float64'>, handle_unknown='error', unknown_value=None, encoded_missing_value=nan, min_frequency=None, max_categories=None)
    Initialize self. See help(type(self)) for accurate signature.
```

```
fit(self, X, y=None)
    Fit the OrdinalEncoder to X.
```

Parameters

X : array-like of shape (n_samples, n_features)
The data to determine the categories of each feature.

y : None
Ignored. This parameter exists only for compatibility with
:class:`~sklearn.pipeline.Pipeline`.

Returns

self : object
Fitted encoder.

```
inverse_transform(self, X)
    Convert the data back to the original representation.
```

Parameters

X : array-like of shape (n_samples, n_encoded_features)
The transformed data.

Returns

X_tr : ndarray of shape (n_samples, n_features)
Inverse transformed array.

```
transform(self, X)
    Transform X to ordinal codes.
```

Parameters

```

-----
X : array-like of shape (n_samples, n_features)
    The data to encode.

Returns
-----
X_out : ndarray of shape (n_samples, n_features)
    Transformed input.

-----
Data and other attributes defined here:

__annotations__ = {'_parameter_constraints': <class 'dict'>}

-----
Methods inherited from sklearn.base.OneToOneFeatureMixin:

get_feature_names_out(self, input_features=None)
    Get output feature names for transformation.

Parameters
-----
input_features : array-like of str or None, default=None
    Input features.

    - If `input_features` is `None`, then `feature_names_in_` is
      used as feature names in. If `feature_names_in_` is not defined,
      then the following input feature names are generated:
      `["x0", "x1", ..., "x(n_features_in_ - 1)"]`.
    - If `input_features` is an array-like, then `input_features` must
      match `feature_names_in_` if `feature_names_in_` is defined.

Returns
-----
feature_names_out : ndarray of str objects
    Same as input features.

-----
Data descriptors inherited from sklearn.base.OneToOneFeatureMixin:

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)

-----
Readonly properties inherited from _BaseEncoder:

infrequent_categories_
    Infrequent categories for each feature.

-----
Methods inherited from sklearn.base.TransformerMixin:

fit_transform(self, X, y=None, **fit_params)
    Fit to data, then transform it.

    Fits transformer to `X` and `y` with optional parameters `fit_params`
    and returns a transformed version of `X`.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Input samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs),
    Target values (None for unsupervised transformations).
    default=None

**fit_params : dict
    Additional fit parameters.

Returns
-----
X_new : ndarray array of shape (n_samples, n_features_new)
    Transformed array.

-----
Methods inherited from sklearn.utils._set_output._SetOutputMixin:

```

```
set_output(self, *, transform=None)
    Set output container.

    See :ref:`sphx_glr_auto_examples_miscellaneous_plot_set_output.py`
    for an example on how to use the API.
```

Parameters

```
transform : {"default", "pandas"}, default=None
    Configure output of `transform` and `fit_transform`.

    - "default": Default output format of a transformer
    - "pandas": DataFrame output
    - None: Transform configuration is unchanged
```

Returns

```
self : estimator instance
    Estimator instance.
```

Class methods inherited from sklearn.utils._set_output._SetOutputMixin:

```
__init_subclass__(auto_wrap_output_keys=('transform',), **kwargs) from builtins.type
    This method is called when a class is subclassed.
```

The default implementation does nothing. It may be overridden to extend subclasses.

Methods inherited from sklearn.base.BaseEstimator:

```
__getstate__(self)
    Helper for pickle.
```

```
__repr__(self, N_CHAR_MAX=700)
    Return repr(self).
```

```
__setstate__(self, state)
```

```
__sklearn_clone__(self)
```

```
get_params(self, deep=True)
    Get parameters for this estimator.
```

Parameters

```
deep : bool, default=True
    If True, will return the parameters for this estimator and
    contained subobjects that are estimators.
```

Returns

```
params : dict
    Parameter names mapped to their values.
```

```
set_params(self, **params)
    Set the parameters of this estimator.
```

The method works on simple estimators as well as on nested objects (such as :class:`~sklearn.pipeline.Pipeline`). The latter have parameters of the form ``<component>__<parameter>`` so that it's possible to update each component of a nested object.

Parameters

```
**params : dict
    Estimator parameters.
```

Returns

```
self : estimator instance
    Estimator instance.
```

Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

```
get_metadata_routing(self)
```



```
|
|   Get metadata routing of this object.
|
|   Please check :ref:`User Guide <metadata_routing>` on how the routing
|   mechanism works.
|
|   Returns
|   -----
|   routing : MetadataRequest
|       A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
|       routing information.
```

```
In [3]: # toy example
import pandas as pd

train_edu = {'educational level': ['Bachelors', 'Masters', 'Bachelors', 'Doctorate', 'HS-grad', 'Masters']}
test_edu = {'educational level': ['HS-grad', 'Masters', 'Masters', 'College', 'Bachelors']}

Xtoy_train = pd.DataFrame(train_edu)
Xtoy_test = pd.DataFrame(test_edu)

# initialize the encoder
cats = [['HS-grad', 'College', 'Bachelors', 'Masters', 'Doctorate']]

enc = OrdinalEncoder(categories = cats) # The ordered list of
# categories need to be provided. By default, the categories are alphabetically ordered!

# fit the training data
enc.fit(Xtoy_train)
# print the categories - not really important because we manually gave the ordered list of categories
print(enc.categories_)
# transform X_train. We could have used enc.fit_transform(X_train) to combine fit and transform
X_train_oe = enc.transform(Xtoy_train)
print(X_train_oe)
# transform X_test
X_test_oe = enc.transform(Xtoy_test) # OrdinalEncoder always throws an error message if
# it encounters an unknown category in test
print(X_test_oe)
```

```
[array(['HS-grad', 'College', 'Bachelors', 'Masters', 'Doctorate'],
      dtype=object)]
```

```
[2.]
[3.]
[2.]
[4.]
[0.]
[3.]]
[[0.]
 [3.]
 [3.]
 [1.]
 [2.]]
```

```
In [4]: # apply OE to the adult dataset
# initialize the encoder
ordinal_fts = ['education'] # if you have more than one ordinal feature, add the feature names here
ordinal_cats = [[' Preschool', ' 1st-4th', ' 5th-6th', ' 7th-8th', ' 9th', ' 10th', ' 11th', ' 12th', ' HS-grad', \
                ' Some-college', ' Assoc-voc', ' Assoc-acdm', ' Bachelors', ' Masters', ' Prof-school', ' Doctorate']]
# ordinal_cats must contain one list per ordinal feature! each list contains the ordered list of categories
# of the corresponding feature

enc = OrdinalEncoder(categories = ordinal_cats) # By default, the categories are alphabetically ordered
# which is NOT what you want usually.

# fit the training data
enc.fit(X_train[ordinal_fts]) # the encoder expects a 2D array, that's why the column name is in a list

# transform X_train. We could use enc.fit_transform(X_train) to combine fit and transform
ordinal_train = enc.transform(X_train[ordinal_fts])
print('transformed train features:')
print(ordinal_train)
# transform X_val
ordinal_val = enc.transform(X_val[ordinal_fts])
print('transformed validation features:')
print(ordinal_val)
# transform X_test
ordinal_test = enc.transform(X_test[ordinal_fts])
print('transformed test features:')
print(ordinal_test)
```

```
transformed train features:
[[10.]
 [ 9.]
 [ 8.]
 ...
 [ 6.]
 [ 8.]
 [12.]]
transformed validation features:
[[14.]
 [13.]
 [ 9.]
 ...
 [12.]
 [ 8.]
 [ 8.]]
transformed test features:
[[12.]
 [ 9.]
 [12.]
 ...
 [ 9.]
 [ 9.]
 [11.]]
```

Unordered categorical data: one-hot encoder

- some categories cannot be ordered. e.g., workclass, relationship status

```
In [5]: from sklearn.preprocessing import OneHotEncoder
        help(OneHotEncoder)
```

Help on class OneHotEncoder in module sklearn.preprocessing._encoders:

```
class OneHotEncoder(BaseEncoder)
| OneHotEncoder(*, categories='auto', drop=None, sparse='deprecated', sparse_output=True, dtype=<class 'numpy.float64'>, handle_unknown='error', min_frequency=None, max_categories=None, feature_name_combiner='concat')
```

Encode categorical features as a one-hot numeric array.

The input to this transformer should be an array-like of integers or strings, denoting the values taken on by categorical (discrete) features. The features are encoded using a one-hot (aka 'one-of-K' or 'dummy') encoding scheme. This creates a binary column for each category and returns a sparse matrix or dense array (depending on the ``sparse_output`` parameter)

By default, the encoder derives the categories based on the unique values in each feature. Alternatively, you can also specify the ``categories`` manually.

This encoding is needed for feeding categorical data to many scikit-learn estimators, notably linear models and SVMs with the standard kernels.

Note: a one-hot encoding of y labels should use a LabelBinarizer instead.

Read more in the :ref:`User Guide <preprocessing_categorical_features>`.

Parameters

categories : 'auto' or a list of array-like, default='auto'

Categories (unique values) per feature:

- 'auto' : Determine categories automatically from the training data.
- list : ``categories[i]`` holds the categories expected in the ith column. The passed categories should not mix strings and numeric values within a single feature, and should be sorted in case of numeric values.

The used categories can be found in the ``categories_`` attribute.

.. versionadded:: 0.20

drop : {'first', 'if_binary'} or an array-like of shape (n_features,), Specifies a methodology to use to drop one of the categories per feature. This is useful in situations where perfectly collinear features cause problems, such as when feeding the resulting data into an unregularized linear regression model.

default=None

However, dropping one category breaks the symmetry of the original representation and can therefore induce a bias in downstream models, for instance for penalized linear classification or regression models.

- None : retain all features (the default).
- 'first' : drop the first category in each feature. If only one category is present, the feature will be dropped entirely.
- 'if_binary' : drop the first category in each feature with two categories. Features with 1 or more than 2 categories are left intact.
- array : ``drop[i]`` is the category in feature ``X[:, i]`` that should be dropped.

When ``max_categories`` or ``min_frequency`` is configured to group infrequent categories, the dropping behavior is handled after the grouping.

.. versionadded:: 0.21
The parameter ``drop`` was added in 0.21.

.. versionchanged:: 0.23
The option ``drop='if_binary'`` was added in 0.23.

.. versionchanged:: 1.1
Support for dropping infrequent categories.

sparse : bool, default=True

Will return sparse matrix if set True else will return an array.

.. deprecated:: 1.2
``sparse`` is deprecated in 1.2 and will be removed in 1.4. Use

```

        `sparse_output` instead.

sparse_output : bool, default=True
    Will return sparse matrix if set True else will return an array.

    .. versionadded:: 1.2
        `sparse` was renamed to `sparse_output`

dtype : number type, default=float
    Desired dtype of output.

handle_unknown : {'error', 'ignore', 'infrequent_if_exist'}, default='error'
    Specifies the way unknown categories are handled during :meth:`transform`.

    - 'error' : Raise an error if an unknown category is present during transform.
    - 'ignore' : When an unknown category is encountered during
      transform, the resulting one-hot encoded columns for this feature
      will be all zeros. In the inverse transform, an unknown category
      will be denoted as None.
    - 'infrequent_if_exist' : When an unknown category is encountered
      during transform, the resulting one-hot encoded columns for this
      feature will map to the infrequent category if it exists. The
      infrequent category will be mapped to the last position in the
      encoding. During inverse transform, an unknown category will be
      mapped to the category denoted `infrequent` if it exists. If the
      `infrequent` category does not exist, then :meth:`transform` and
      :meth:`inverse_transform` will handle an unknown category as with
      `handle_unknown='ignore'`. Infrequent categories exist based on
      `min_frequency` and `max_categories`. Read more in the
      :ref:`User Guide <encoder_infrequent_categories>`.

    .. versionchanged:: 1.1
        `infrequent_if_exist` was added to automatically handle unknown
        categories and infrequent categories.

min_frequency : int or float, default=None
    Specifies the minimum frequency below which a category will be
    considered infrequent.

    - If `int`, categories with a smaller cardinality will be considered
      infrequent.

    - If `float`, categories with a smaller cardinality than
      `min_frequency * n_samples` will be considered infrequent.

    .. versionadded:: 1.1
        Read more in the :ref:`User Guide <encoder_infrequent_categories>`.

max_categories : int, default=None
    Specifies an upper limit to the number of output features for each input
    feature when considering infrequent categories. If there are infrequent
    categories, `max_categories` includes the category representing the
    infrequent categories along with the frequent categories. If `None`,
    there is no limit to the number of output features.

    .. versionadded:: 1.1
        Read more in the :ref:`User Guide <encoder_infrequent_categories>`.

feature_name_combiner : "concat" or callable, default="concat"
    Callable with signature `def callable(input_feature, category)` that returns a
    string. This is used to create feature names to be returned by
    :meth:`get_feature_names_out`.

    `concat` concatenates encoded feature name and category with
    `feature + "_" + str(category)`. E.g. feature X with values 1, 6, 7 create
    feature names `X_1, X_6, X_7`.

    .. versionadded:: 1.3

Attributes
-----
categories_ : list of arrays
    The categories of each feature determined during fitting
    (in order of the features in X and corresponding with the output
    of `transform`). This includes the category specified in `drop`
    (if any).

drop_idx_ : array of shape (n_features,)
    - `drop_idx[i]` is the index in `categories_[i]` of the category

```

```

        to be dropped for each feature.
    - ``drop_idx[i] = None`` if no category is to be dropped from the
      feature with index ``i``, e.g. when ``drop='if_binary'`` and the
      feature isn't binary.
    - ``drop_idx_ = None`` if all the transformed features will be
      retained.

    If infrequent categories are enabled by setting ``min_frequency`` or
    ``max_categories`` to a non-default value and ``drop_idx[i]`` corresponds
    to a infrequent category, then the entire infrequent category is
    dropped.

    .. versionchanged:: 0.23
        Added the possibility to contain ``None`` values.

    infrequent_categories_ : list of ndarray
        Defined only if infrequent categories are enabled by setting
        ``min_frequency`` or ``max_categories`` to a non-default value.
        ``infrequent_categories_[i]`` are the infrequent categories for feature
        ``i``. If the feature ``i`` has no infrequent categories
        ``infrequent_categories_[i]`` is None.

    .. versionadded:: 1.1

    n_features_in_ : int
        Number of features seen during :term:`fit`.

    .. versionadded:: 1.0

    feature_names_in_ : ndarray of shape (`n_features_in`,)
        Names of features seen during :term:`fit`. Defined only when `X`
        has feature names that are all strings.

    .. versionadded:: 1.0

    feature_name_combiner : callable or None
        Callable with signature ``def callable(input_feature, category)`` that returns a
        string. This is used to create feature names to be returned by
        :meth:`get_feature_names_out`.

    .. versionadded:: 1.3

```

See Also

```

OrdinalEncoder : Performs an ordinal (integer)
                  encoding of the categorical features.
TargetEncoder : Encodes categorical features using the target.
sklearn.feature_extraction.DictVectorizer : Performs a one-hot encoding of
      dictionary items (also handles string-valued features).
sklearn.feature_extraction.FeatureHasher : Performs an approximate one-hot
      encoding of dictionary items or strings.
LabelBinarizer : Binarizes labels in a one-vs-all
                  fashion.
MultiLabelBinarizer : Transforms between iterable of
      iterables and a multilabel format, e.g. a (samples x classes) binary
      matrix indicating the presence of a class label.

```

Examples

```

Given a dataset with two features, we let the encoder find the unique
values per feature and transform the data to a binary one-hot encoding.

>>> from sklearn.preprocessing import OneHotEncoder

One can discard categories not seen during `fit`:

>>> enc = OneHotEncoder(handle_unknown='ignore')
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
>>> enc.fit(X)
OneHotEncoder(handle_unknown='ignore')
>>> enc.categories_
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> enc.transform([['Female', 1], ['Male', 4]]).toarray()
array([[1., 0., 1., 0., 0.],
       [0., 1., 0., 0., 0.]])
>>> enc.inverse_transform([0, 1, 1, 0, 0], [0, 0, 0, 1, 0])
array([['Male', 1],
       [None, 2]], dtype=object)
>>> enc.get_feature_names_out(['gender', 'group'])

```

```

array(['gender_Female', 'gender_Male', 'group_1', 'group_2', 'group_3'], ...)

One can always drop the first column for each feature:

>>> drop_enc = OneHotEncoder(drop='first').fit(X)
>>> drop_enc.categories_
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> drop_enc.transform([['Female', 1], ['Male', 2]]).toarray()
array([[0., 0., 0.],
       [1., 1., 0.]])

Or drop a column for feature only having 2 categories:

>>> drop_binary_enc = OneHotEncoder(drop='if_binary').fit(X)
>>> drop_binary_enc.transform([['Female', 1], ['Male', 2]]).toarray()
array([[0., 1., 0., 0.],
       [1., 0., 1., 0.]])

One can change the way feature names are created.

>>> def custom_combiner(feature, category):
...     return str(feature) + "_" + type(category).__name__ + "_" + str(category)
>>> custom_fnames_enc = OneHotEncoder(feature_name_combiner=custom_combiner).fit(X)
>>> custom_fnames_enc.get_feature_names_out()
array(['x0_str_Female', 'x0_str_Male', 'x1_int_1', 'x1_int_2', 'x1_int_3'],
      dtype=object)

Infrequent categories are enabled by setting `max_categories` or `min_frequency`.

>>> import numpy as np
>>> X = np.array([["a"] * 5 + ["b"] * 20 + ["c"] * 10 + ["d"] * 3], dtype=object).T
>>> ohe = OneHotEncoder(max_categories=3, sparse_output=False).fit(X)
>>> ohe.infrequent_categories_
[array(['a', 'd'], dtype=object)]
>>> ohe.transform([["a"], ["b"]])
array([[0., 0., 1.],
       [1., 0., 0.]])

Method resolution order:
  OneHotEncoder
  _BaseEncoder
  sklearn.base.TransformerMixin
  sklearn.utils._set_output._SetOutputMixin
  sklearn.base.BaseEstimator
  sklearn.utils._metadata_requests._MetadataRequester
  builtins.object

Methods defined here:

  __init__(self, *, categories='auto', drop=None, sparse='deprecated', sparse_output=True, dtype=<class 'numpy.f
loat64'>, handle_unknown='error', min_frequency=None, max_categories=None, feature_name_combiner='concat')
    Initialize self. See help(type(self)) for accurate signature.

  fit(self, X, y=None)
    Fit OneHotEncoder to X.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        The data to determine the categories of each feature.

    y : None
        Ignored. This parameter exists only for compatibility with
        :class:`~sklearn.pipeline.Pipeline`.

    Returns
    -----
    self
        Fitted encoder.

  get_feature_names_out(self, input_features=None)
    Get output feature names for transformation.

    Parameters
    -----
    input_features : array-like of str or None, default=None
        Input features.

        - If `input_features` is `None`, then `feature_names_in_` is

```

```

        used as feature names in. If `feature_names_in_` is not defined,
        then the following input feature names are generated:
        `["x0", "x1", ..., "x(n_features_in_ - 1)"]`.
    - If `input_features` is an array-like, then `input_features` must
      match `feature_names_in_` if `feature_names_in_` is defined.

Returns
-----
feature_names_out : ndarray of str objects
    Transformed feature names.

inverse_transform(self, X)
    Convert the data back to the original representation.

    When unknown categories are encountered (all zeros in the
    one-hot encoding), `None` is used to represent this category. If the
    feature with the unknown category has a dropped category, the dropped
    category will be its inverse.

    For a given input feature, if there is an infrequent category,
    'infrequent_sklearn' will be used to represent the infrequent category.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_encoded_features)
    The transformed data.

Returns
-----
X_tr : ndarray of shape (n_samples, n_features)
    Inverse transformed array.

transform(self, X)
    Transform X using one-hot encoding.

    If there are infrequent categories for a feature, the infrequent
    categories will be grouped into a single category.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    The data to encode.

Returns
-----
X_out : {ndarray, sparse matrix} of shape (n_samples, n_encoded_features)
    Transformed input. If `sparse_output=True`, a sparse matrix will be
    returned.

```

```

Data and other attributes defined here:

__annotations__ = {'_parameter_constraints': <class 'dict'>}

```

```

Readonly properties inherited from _BaseEncoder:

infrequent_categories_
    Infrequent categories for each feature.

```

```

Methods inherited from sklearn.base.TransformerMixin:

fit_transform(self, X, y=None, **fit_params)
    Fit to data, then transform it.

    Fits transformer to `X` and `y` with optional parameters `fit_params`
    and returns a transformed version of `X`.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Input samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs),
    Target values (None for unsupervised transformations).
    default=None

**fit_params : dict
    Additional fit parameters.

```

Returns

X_new : ndarray array of shape (n_samples, n_features_new)
Transformed array.

Methods inherited from sklearn.utils._set_output._SetOutputMixin:

set_output(self, *, transform=None)
Set output container.

See :ref:`sphx_glr_auto_examples_miscellaneous_plot_set_output.py`
for an example on how to use the API.

Parameters

transform : {"default", "pandas"}, default=None
Configure output of `transform` and `fit_transform`.

- "default": Default output format of a transformer
- "pandas": DataFrame output
- None: Transform configuration is unchanged

Returns

self : estimator instance
Estimator instance.

Class methods inherited from sklearn.utils._set_output._SetOutputMixin:

__init_subclass__(auto_wrap_output_keys=('transform',), **kwargs) from builtins.type
Set the ``set_{method}_request`` methods.

This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It looks for the information available in the set default values which are set using ``__metadata_request__*`` class attributes, or inferred from method signatures.

The ``__metadata_request__*`` class attributes are used when a method does not explicitly accept a metadata through its arguments or if the developer would like to specify a request value for those metadata which are different from the default ``None``.

References

.. [1] <https://www.python.org/dev/peps/pep-0487>

Data descriptors inherited from sklearn.utils._set_output._SetOutputMixin:

__dict__
dictionary for instance variables (if defined)

__weakref__
list of weak references to the object (if defined)

Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)
Helper for pickle.

__repr__(self, N_CHAR_MAX=700)
Return repr(self).

__setstate__(self, state)

__sklearn_clone__(self)

get_params(self, deep=True)
Get parameters for this estimator.

Parameters

deep : bool, default=True
If True, will return the parameters for this estimator and contained subobjects that are estimators.


```

Returns
-----
params : dict
    Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

    The method works on simple estimators as well as on nested objects
    (such as :class:`~sklearn.pipeline.Pipeline`). The latter have
    parameters of the form ``<component>__<parameter>`` so that it's
    possible to update each component of a nested object.

Parameters
-----
**params : dict
    Estimator parameters.

Returns
-----
self : estimator instance
    Estimator instance.

```

```

Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

get_metadata_routing(self)
    Get metadata routing of this object.

    Please check :ref:`User Guide <metadata_routing>` on how the routing
    mechanism works.

Returns
-----
routing : MetadataRequest
    A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
    routing information.

```

```

In [7]: # toy example
train = {'gender':['Male','Female','Unknown','Male','Female','Female'],\
        'browser':['Safari','Safari','Internet Explorer','Chrome','Chrome','Internet Explorer']}
test = {'gender':['Female','Male','Unknown','Female'], 'browser':['Chrome','Firefox','Internet Explorer','Safari']}

Xtoy_train = pd.DataFrame(train)
Xtoy_test = pd.DataFrame(test)

ftrs = ['gender','browser']

# initialize the encoder
enc = OneHotEncoder(sparse=False) # by default, OneHotEncoder returns a sparse matrix. sparse=False returns a 2D
# fit the training data
enc.fit(Xtoy_train)
print('categories:',enc.categories_)
print('feature names:',enc.get_feature_names_out(ftrs))
# transform X_train
X_train_ohe = enc.transform(Xtoy_train)
#print(X_train_ohe)
# do all of this in one step
X_train_ohe = enc.fit_transform(Xtoy_train)
print('X_train transformed')
print(X_train_ohe)

# transform X_test
X_test_ohe = enc.transform(Xtoy_test)
print('X_test transformed')
print(X_test_ohe)

```

```

categories: [array(['Female', 'Male', 'Unknown'], dtype=object), array(['Chrome', 'Internet Explorer', 'Safari'],
dtype=object)]
feature names: ['gender_Female' 'gender_Male' 'gender_Unknown' 'browser_Chrome'
'browser_Internet Explorer' 'browser_Safari']
X_train transformed
[[0. 1. 0. 0. 0. 1.]
 [1. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 1. 0.]
 [0. 1. 0. 1. 0. 0.]
 [1. 0. 0. 1. 0. 0.]
 [1. 0. 0. 0. 1. 0.]]

```

/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

ValueError Traceback (most recent call last)

Cell In[7], line 26

```

23 print(X_train_ohe)
25 # transform X_test
--> 26 X_test_ohe = enc.transform(Xtoy_test)
27 print('X_test transformed')
28 print(X_test_ohe)

```

File ~/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/utils/_set_output.py:140, in _wrap_method_output.<locals>.wrapped(self, X, *args, **kwargs)

```

138 @wraps(f)
139 def wrapped(self, X, *args, **kwargs):
--> 140     data_to_wrap = f(self, X, *args, **kwargs)
141     if isinstance(data_to_wrap, tuple):
142         # only wrap the first output for cross decomposition
143         return_tuple = (
144             _wrap_data_with_container(method, data_to_wrap[0], X, self),
145             *data_to_wrap[1:],
146         )

```

File ~/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:1016, in OneHotEncoder.transform(self, X)

```

1011 # validation of X happens in _check_X called by _transform
1012 warn_on_unknown = self.drop is not None and self.handle_unknown in {
1013     "ignore",
1014     "infrequent_if_exist",
1015 }
-> 1016 X_int, X_mask = self._transform(
1017     X,
1018     handle_unknown=self.handle_unknown,
1019     force_all_finite="allow-nan",
1020     warn_on_unknown=warn_on_unknown,
1021 )
1023 n_samples, n_features = X_int.shape
1025 if self._drop_idx_after_grouping is not None:

```

File ~/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:199, in _BaseEncoder._transform(self, X, handle_unknown, force_all_finite, warn_on_unknown, ignore_category_indices)

```

194 if handle_unknown == "error":
195     msg = (
196         "Found unknown categories {0} in column {1}"
197         " during transform".format(diff, i)
198     )
-> 199     raise ValueError(msg)
200 else:
201     if warn_on_unknown:

```

ValueError: Found unknown categories ['Firefox'] in column 1 during transform

```

In [9]: # apply OHE to the adult dataset

# let's collect all categorical features first
onehot_fts = ['workclass', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country']
# initialize the encoder
enc = OneHotEncoder(sparse=False, handle_unknown='ignore') # by default, OneHotEncoder returns a sparse matrix. sp
# fit the training data
enc.fit(X_train[onehot_fts])
print('feature names:', enc.get_feature_names_out(onehot_fts))

```

```

feature names: ['workclass_?' 'workclass_ Federal-gov' 'workclass_ Local-gov'
'workclass_ Never-worked' 'workclass_ Private' 'workclass_ Self-emp-inc'
'workclass_ Self-emp-not-inc' 'workclass_ State-gov'
'workclass_ Without-pay' 'marital-status_ Divorced'
'marital-status_ Married-AF-spouse' 'marital-status_ Married-civ-spouse'
'marital-status_ Married-spouse-absent' 'marital-status_ Never-married'
'marital-status_ Separated' 'marital-status_ Widowed' 'occupation_?'
'occupation_ Adm-clerical' 'occupation_ Armed-Forces'
'occupation_ Craft-repair' 'occupation_ Exec-managerial'
'occupation_ Farming-fishing' 'occupation_ Handlers-cleaners'
'occupation_ Machine-op-inspct' 'occupation_ Other-service'
'occupation_ Priv-house-serv' 'occupation_ Prof-specialty'
'occupation_ Protective-serv' 'occupation_ Sales'
'occupation_ Tech-support' 'occupation_ Transport-moving'
'relationship_ Husband' 'relationship_ Not-in-family'
'relationship_ Other-relative' 'relationship_ Own-child'
'relationship_ Unmarried' 'relationship_ Wife' 'race_ Amer-Indian-Eskimo'
'race_ Asian-Pac-Islander' 'race_ Black' 'race_ Other' 'race_ White'
'sex_ Female' 'sex_ Male' 'native-country_?' 'native-country_ Cambodia'
'native-country_ Canada' 'native-country_ China'
'native-country_ Columbia' 'native-country_ Cuba'
'native-country_ Dominican-Republic' 'native-country_ Ecuador'
'native-country_ El-Salvador' 'native-country_ England'
'native-country_ France' 'native-country_ Germany'
'native-country_ Greece' 'native-country_ Guatemala'
'native-country_ Haiti' 'native-country_ Holand-Netherlands'
'native-country_ Honduras' 'native-country_ Hong'
'native-country_ Hungary' 'native-country_ India' 'native-country_ Iran'
'native-country_ Ireland' 'native-country_ Italy'
'native-country_ Jamaica' 'native-country_ Japan' 'native-country_ Laos'
'native-country_ Mexico' 'native-country_ Nicaragua'
'native-country_ Outlying-US(Guam-USVI-etc)' 'native-country_ Peru'
'native-country_ Philippines' 'native-country_ Poland'
'native-country_ Portugal' 'native-country_ Puerto-Rico'
'native-country_ Scotland' 'native-country_ South'
'native-country_ Taiwan' 'native-country_ Thailand'
'native-country_ Trinidad&Tobago' 'native-country_ United-States'
'native-country_ Vietnam' 'native-country_ Yugoslavia']

```

```

/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
  warnings.warn(

```

```

In [10]: # transform X_train
onehot_train = enc.transform(X_train[onehot_fts])
print('transformed train features:')
print(onehot_train)
# transform X_val
onehot_val = enc.transform(X_val[onehot_fts])
print('transformed val features:')
print(onehot_val)
# transform X_test
onehot_test = enc.transform(X_test[onehot_fts])
print('transformed test features:')
print(onehot_test)

```

```

transformed train features:
[[0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 ...
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]
transformed val features:
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]
transformed test features:
[[0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]

```

Quiz 1

Please explain how you would encode the race feature below and what would be the output of the encoder. Do not write code. The goal of this quiz is to test your conceptual understanding so write text and the output array.

```
race = ['Amer-Indian-Eskimo', 'White', 'Black', 'Asian-Pac-Islander', 'Black', 'White', 'White']
```

By the end of this lecture, you will be able to

- apply one-hot encoding or ordinal encoding to categorical variables
- **apply scaling and normalization to continuous variables**

Continuous features: MinMaxScaler

- If the continuous feature values are reasonably bounded, MinMaxScaler is a good way to scale the features.
- Age is expected to be within the range of 0 and 100.
- Number of hours worked per week is in the range of 0 to 80.
- If unsure, plot the histogram of the feature to verify or just go with the standard scaler!

```
In [11]: from sklearn.preprocessing import MinMaxScaler
         help(MinMaxScaler)
```

Help on class MinMaxScaler in module sklearn.preprocessing._data:

```
class MinMaxScaler(sklearn.base.OneToOneFeatureMixin, sklearn.base.TransformerMixin, sklearn.base.BaseEstimator)
|   MinMaxScaler(feature_range=(0, 1), *, copy=True, clip=False)
|
|   Transform features by scaling each feature to a given range.
|
|   This estimator scales and translates each feature individually such
|   that it is in the given range on the training set, e.g. between
|   zero and one.
|
|   The transformation is given by::
|
|       
$$X_{std} = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))$$

|       
$$X_{scaled} = X_{std} * (max - min) + min$$

|
|   where min, max = feature_range.
|
|   This transformation is often used as an alternative to zero mean,
|   unit variance scaling.
|
|   Read more in the :ref:`User Guide <preprocessing_scaler>`.
|
|   Parameters
|   -----
|   feature_range : tuple (min, max), default=(0, 1)
|       Desired range of transformed data.
|
|   copy : bool, default=True
|       Set to False to perform inplace row normalization and avoid a
|       copy (if the input is already a numpy array).
|
|   clip : bool, default=False
|       Set to True to clip transformed values of held-out data to
|       provided `feature range`.
|
|   .. versionadded:: 0.24
|
|   Attributes
|   -----
|   min_ : ndarray of shape (n_features,)
|       Per feature adjustment for minimum. Equivalent to
|       ``min - X.min(axis=0) * self.scale_``
|
|   scale_ : ndarray of shape (n_features,)
|       Per feature relative scaling of the data. Equivalent to
|       ``(max - min) / (X.max(axis=0) - X.min(axis=0))``
|
|   .. versionadded:: 0.17
|       *scale_* attribute.
|
|   data_min_ : ndarray of shape (n_features,)
|       Per feature minimum seen in the data
|
|   .. versionadded:: 0.17
|       *data_min_*
|
|   data_max_ : ndarray of shape (n_features,)
|       Per feature maximum seen in the data
|
|   .. versionadded:: 0.17
|       *data_max_*
|
|   data_range_ : ndarray of shape (n_features,)
|       Per feature range ``(data_max_ - data_min_)`` seen in the data
|
|   .. versionadded:: 0.17
|       *data_range_*
|
|   n_features_in_ : int
|       Number of features seen during :term:`fit`.
|
|   .. versionadded:: 0.24
|
|   n_samples_seen_ : int
|       The number of samples processed by the estimator.
|       It will be reset on new calls to fit, but increments across
|       ``partial_fit`` calls.
```

feature_names_in_ : ndarray of shape (``n_features_in_``,)
Names of features seen during `:term:`fit``. Defined only when ``X``
has feature names that are all strings.

.. versionadded:: 1.0

See Also

`minmax_scale` : Equivalent function without the estimator API.

Notes

NaNs are treated as missing values: disregarded in fit, and maintained in transform.

For a comparison of the different scalers, transformers, and normalizers,
see `:ref:`examples/preprocessing/plot_all_scaling.py``
`<sphinx_glr_auto_examples_preprocessing_plot_all_scaling.py>`.

Examples

```
>>> from sklearn.preprocessing import MinMaxScaler
>>> data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]
>>> scaler = MinMaxScaler()
>>> print(scaler.fit(data))
MinMaxScaler()
>>> print(scaler.data_max_)
[ 1. 18.]
>>> print(scaler.transform(data))
[[0.  0. ]
 [0.25 0.25]
 [0.5  0.5 ]
 [1.  1.  ]]
>>> print(scaler.transform([[2, 2]]))
[[1.5 0.  ]]
```

Method resolution order:

MinMaxScaler
sklearn.base.OneToOneFeatureMixin
sklearn.base.TransformerMixin
sklearn.utils._set_output._SetOutputMixin
sklearn.base.BaseEstimator
sklearn.utils.metadata_requests._MetadataRequester
builtins.object

Methods defined here:

`__init__(self, feature_range=(0, 1), *, copy=True, clip=False)`
Initialize self. See `help(type(self))` for accurate signature.

`fit(self, X, y=None)`
Compute the minimum and maximum to be used for later scaling.

Parameters

`X` : array-like of shape (n_samples, n_features)
The data used to compute the per-feature minimum and maximum
used for later scaling along the features axis.

`y` : None
Ignored.

Returns

`self` : object
Fitted scaler.

`inverse_transform(self, X)`
Undo the scaling of `X` according to `feature_range`.

Parameters

`X` : array-like of shape (n_samples, n_features)
Input data that will be transformed. It cannot be sparse.

Returns

`Xt` : ndarray of shape (n_samples, n_features)
Transformed data.

```
partial_fit(self, X, y=None)
    Online computation of min and max on X for later scaling.

    All of X is processed as a single batch. This is intended for cases
    when :meth:`fit` is not feasible due to very large number of
    `n_samples` or because X is read from a continuous stream.
```

Parameters

X : array-like of shape (n_samples, n_features)
The data used to compute the mean and standard deviation
used for later scaling along the features axis.

y : None
Ignored.

Returns

self : object
Fitted scaler.

```
transform(self, X)
    Scale features of X according to feature_range.
```

Parameters

X : array-like of shape (n_samples, n_features)
Input data that will be transformed.

Returns

Xt : ndarray of shape (n_samples, n_features)
Transformed data.

Data and other attributes defined here:

```
__annotations__ = {'_parameter_constraints': <class 'dict'>}
```

Methods inherited from sklearn.base.OneToOneFeatureMixin:

```
get_feature_names_out(self, input_features=None)
    Get output feature names for transformation.
```

Parameters

input_features : array-like of str or None, default=None
Input features.

- If `input_features` is `None`, then `feature_names_in_` is used as feature names in. If `feature_names_in_` is not defined, then the following input feature names are generated: `["x0", "x1", ..., "x(n_features_in_ - 1)"]`.
- If `input_features` is an array-like, then `input_features` must match `feature_names_in_` if `feature_names_in_` is defined.

Returns

feature_names_out : ndarray of str objects
Same as input features.

Data descriptors inherited from sklearn.base.OneToOneFeatureMixin:

```
__dict__
    dictionary for instance variables (if defined)
```

```
__weakref__
    list of weak references to the object (if defined)
```

Methods inherited from sklearn.base.TransformerMixin:

```
fit_transform(self, X, y=None, **fit_params)
    Fit to data, then transform it.
```

Fits transformer to `X` and `y` with optional parameters `fit_params`

```

    and returns a transformed version of `X`.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Input samples.

    y : array-like of shape (n_samples,) or (n_samples, n_outputs),
        Target values (None for unsupervised transformations).
        default=None

    **fit_params : dict
        Additional fit parameters.

    Returns
    -----
    X_new : ndarray array of shape (n_samples, n_features_new)
        Transformed array.

```

```

Methods inherited from sklearn.utils._set_output._SetOutputMixin:

set_output(self, *, transform=None)
    Set output container.

    See :ref:`sphx_glr_auto_examples_misellaneous_plot_set_output.py`
    for an example on how to use the API.

    Parameters
    -----
    transform : {"default", "pandas"}, default=None
        Configure output of `transform` and `fit_transform`.

        - "default": Default output format of a transformer
        - "pandas": DataFrame output
        - None: Transform configuration is unchanged

    Returns
    -----
    self : estimator instance
        Estimator instance.

```

```

Class methods inherited from sklearn.utils._set_output._SetOutputMixin:

__init_subclass__(auto_wrap_output_keys=('transform',), **kwargs) from builtins.type
    This method is called when a class is subclassed.

    The default implementation does nothing. It may be
    overridden to extend subclasses.

```

```

Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)
    Helper for pickle.

__repr__(self, N_CHAR_MAX=700)
    Return repr(self).

__setstate__(self, state)

__sklearn_clone__(self)

get_params(self, deep=True)
    Get parameters for this estimator.

    Parameters
    -----
    deep : bool, default=True
        If True, will return the parameters for this estimator and
        contained subobjects that are estimators.

    Returns
    -----
    params : dict
        Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

```



```

|
| The method works on simple estimators as well as on nested objects
| (such as :class:`~sklearn.pipeline.Pipeline`). The latter have
| parameters of the form ``<component>__<parameter>`` so that it's
| possible to update each component of a nested object.
|
| Parameters
| -----
| **params : dict
|     Estimator parameters.
|
| Returns
| -----
| self : estimator instance
|     Estimator instance.
|
| -----
| Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:
|
| get_metadata_routing(self)
|     Get metadata routing of this object.
|
| Please check :ref:`User Guide <metadata_routing>` on how the routing
| mechanism works.
|
| Returns
| -----
| routing : MetadataRequest
|     A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
|     routing information.

```

```

In [12]: # toy data
# let's assume we have two continuous features:
train = {'age': [32, 65, 13, 68, 42, 75, 32], 'number of hours worked': [0, 40, 10, 60, 40, 20, 40]}
test = {'age': [83, 26, 10, 60], 'number of hours worked': [0, 40, 0, 60]}

# (value - min) / (max - min), if value is 32, min is 13 and max is 75, then we have 19 / 62 = 0.3064

Xtoy_train = pd.DataFrame(train)
Xtoy_test = pd.DataFrame(test)

scaler = MinMaxScaler()
scaler.fit(Xtoy_train)
print(scaler.transform(Xtoy_train))
print(scaler.transform(Xtoy_test)) # note how scaled X_test contains values larger than 1 and smaller than 0.

[[0.30645161 0.        ]
 [0.83870968 0.66666667]
 [0.        0.16666667]
 [0.88709677 1.        ]
 [0.46774194 0.66666667]
 [1.        0.33333333]
 [0.30645161 0.66666667]]
[[ 1.12903226 0.        ]
 [ 0.20967742 0.66666667]
 [-0.0483871  0.        ]
 [ 0.75806452 1.        ]]

```

```

In [13]: # adult data

minmax_fts = ['age', 'hours-per-week']

scaler = MinMaxScaler()
scaler.fit(X_train[minmax_fts])
print(scaler.transform(X_train[minmax_fts]))
print(scaler.transform(X_val[minmax_fts]))
print(scaler.transform(X_test[minmax_fts]))

```

```

[[0.19178082 0.39795918]
 [0.32876712 0.39795918]
 [0.60273973 0.5          ]
 ...
 [0.01369863 0.19387755]
 [0.45205479 0.84693878]
 [0.23287671 0.60204082]]
[[0.35616438 0.5          ]
 [0.68493151 0.39795918]
 [0.09589041 0.39795918]
 ...
 [0.09589041 0.19387755]
 [0.02739726 0.44897959]
 [0.38356164 0.39795918]]
[[0.06849315 0.39795918]
 [0.23287671 0.39795918]
 [0.43835616 0.5          ]
 ...
 [0.20547945 0.39795918]
 [0.21917808 0.37755102]
 [0.08219178 0.35714286]]

```

Continuous features: StandardScaler

- If the continuous feature values follow a tailed distribution, StandardScaler is better to use!
- Salaries are a good example. Most people earn less than 100k but there are a small number of super-rich people.

```

In [14]: from sklearn.preprocessing import StandardScaler
         help(StandardScaler)

```

Help on class StandardScaler in module sklearn.preprocessing._data:

```
class StandardScaler(sklearn.base.OneToOneFeatureMixin, sklearn.base.TransformerMixin, sklearn.base.BaseEstimator)
| StandardScaler(*, copy=True, with_mean=True, with_std=True)
|
| Standardize features by removing the mean and scaling to unit variance.
|
| The standard score of a sample `x` is calculated as:
|
|     
$$z = (x - u) / s$$

|
| where `u` is the mean of the training samples or zero if `with_mean=False`,
| and `s` is the standard deviation of the training samples or one if
| `with_std=False`.
|
| Centering and scaling happen independently on each feature by computing
| the relevant statistics on the samples in the training set. Mean and
| standard deviation are then stored to be used on later data using
| :meth:`transform`.
|
| Standardization of a dataset is a common requirement for many
| machine learning estimators: they might behave badly if the
| individual features do not more or less look like standard normally
| distributed data (e.g. Gaussian with 0 mean and unit variance).
|
| For instance many elements used in the objective function of
| a learning algorithm (such as the RBF kernel of Support Vector
| Machines or the L1 and L2 regularizers of linear models) assume that
| all features are centered around 0 and have variance in the same
| order. If a feature has a variance that is orders of magnitude larger
| than others, it might dominate the objective function and make the
| estimator unable to learn from other features correctly as expected.
|
| This scaler can also be applied to sparse CSR or CSC matrices by passing
| `with_mean=False` to avoid breaking the sparsity structure of the data.
|
| Read more in the :ref:`User Guide <preprocessing_scaler>`.
|
| Parameters
| -----
| copy : bool, default=True
|     If False, try to avoid a copy and do inplace scaling instead.
|     This is not guaranteed to always work inplace; e.g. if the data is
|     not a NumPy array or scipy.sparse CSR matrix, a copy may still be
|     returned.
|
| with_mean : bool, default=True
|     If True, center the data before scaling.
|     This does not work (and will raise an exception) when attempted on
|     sparse matrices, because centering them entails building a dense
|     matrix which in common use cases is likely to be too large to fit in
|     memory.
|
| with_std : bool, default=True
|     If True, scale the data to unit variance (or equivalently,
|     unit standard deviation).
|
| Attributes
| -----
| scale_ : ndarray of shape (n_features,) or None
|     Per feature relative scaling of the data to achieve zero mean and unit
|     variance. Generally this is calculated using `np.sqrt(var_)`. If a
|     variance is zero, we can't achieve unit variance, and the data is left
|     as-is, giving a scaling factor of 1. `scale_` is equal to `None`
|     when `with_std=False`.
|
|     .. versionadded:: 0.17
|        *scale_*
|
| mean_ : ndarray of shape (n_features,) or None
|     The mean value for each feature in the training set.
|     Equal to ``None`` when ``with_mean=False``.
|
| var_ : ndarray of shape (n_features,) or None
|     The variance for each feature in the training set. Used to compute
|     `scale_`. Equal to ``None`` when ``with_std=False``.
|
| n_features_in_ : int
|     Number of features seen during :term:`fit`.
```

```

.. versionadded:: 0.24

feature_names_in_ : ndarray of shape (n_features_in_,)
    Names of features seen during :term:`fit`. Defined only when `X`
    has feature names that are all strings.

.. versionadded:: 1.0

n_samples_seen_ : int or ndarray of shape (n_features,)
    The number of samples processed by the estimator for each feature.
    If there are no missing samples, the ``n_samples_seen`` will be an
    integer, otherwise it will be an array of dtype int. If
    ``sample_weights`` are used it will be a float (if no missing data)
    or an array of dtype float that sums the weights seen so far.
    Will be reset on new calls to fit, but increments across
    ``partial_fit`` calls.

See Also
-----
scale : Equivalent function without the estimator API.

:class:`~sklearn.decomposition.PCA` : Further removes the linear
    correlation across features with 'whiten=True'.

Notes
-----
NaNs are treated as missing values: disregarded in fit, and maintained in
transform.

We use a biased estimator for the standard deviation, equivalent to
`numpy.std(x, ddof=0)`. Note that the choice of `ddof` is unlikely to
affect model performance.

For a comparison of the different scalers, transformers, and normalizers,
see :ref:`examples/preprocessing/plot_all_scaling.py`
<sphinx_glr_auto_examples_preprocessing_plot_all_scaling.py>`.

Examples
-----
>>> from sklearn.preprocessing import StandardScaler
>>> data = [[0, 0], [0, 0], [1, 1], [1, 1]]
>>> scaler = StandardScaler()
>>> print(scaler.fit(data))
StandardScaler()
>>> print(scaler.mean_)
[0.5 0.5]
>>> print(scaler.transform(data))
[[-1. -1.]
 [-1. -1.]
 [ 1.  1.]
 [ 1.  1.]]
>>> print(scaler.transform([[2, 2]]))
[[3. 3.]]

Method resolution order:
    StandardScaler
    sklearn.base.OneToOneFeatureMixin
    sklearn.base.TransformerMixin
    sklearn.utils._set_output._SetOutputMixin
    sklearn.base.BaseEstimator
    sklearn.utils._metadata_requests._MetadataRequester
    builtins.object

Methods defined here:

__init__(self, *, copy=True, with_mean=True, with_std=True)
    Initialize self. See help(type(self)) for accurate signature.

fit(self, X, y=None, sample_weight=None)
    Compute the mean and std to be used for later scaling.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    The data used to compute the mean and standard deviation
    used for later scaling along the features axis.

y : None

```

```

        Ignored.

sample_weight : array-like of shape (n_samples,), default=None
    Individual weights for each sample.

    .. versionadded:: 0.24
        parameter *sample_weight* support to StandardScaler.

Returns
-----
self : object
    Fitted scaler.

inverse_transform(self, X, copy=None)
    Scale back the data to the original representation.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    The data used to scale along the features axis.
copy : bool, default=None
    Copy the input X or not.

Returns
-----
X_tr : {ndarray, sparse matrix} of shape (n_samples, n_features)
    Transformed array.

partial_fit(self, X, y=None, sample_weight=None)
    Online computation of mean and std on X for later scaling.

    All of X is processed as a single batch. This is intended for cases
    when :meth:`fit` is not feasible due to very large number of
    `n_samples` or because X is read from a continuous stream.

    The algorithm for incremental mean and std is given in Equation 1.5a,b
    in Chan, Tony F., Gene H. Golub, and Randall J. LeVeque. "Algorithms
    for computing the sample variance: Analysis and recommendations."
    The American Statistician 37.3 (1983): 242-247.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    The data used to compute the mean and standard deviation
    used for later scaling along the features axis.

y : None
    Ignored.

sample_weight : array-like of shape (n_samples,), default=None
    Individual weights for each sample.

    .. versionadded:: 0.24
        parameter *sample_weight* support to StandardScaler.

Returns
-----
self : object
    Fitted scaler.

set_fit_request(self: sklearn.preprocessing._data.StandardScaler, *, sample_weight: Union[bool, NoneType, str]
= '$UNCHANGED$') -> sklearn.preprocessing._data.StandardScaler
    Request metadata passed to the ``fit`` method.

    Note that this method is only relevant if
    ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
    Please see :ref:`User Guide <metadata_routing>` on how the routing
    mechanism works.

    The options for each parameter are:

    - ``True``: metadata is requested, and passed to ``fit`` if provided. The request is ignored if metadata i
s not provided.

    - ``False``: metadata is not requested and the meta-estimator will not pass it to ``fit``.

    - ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.

    - ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original n

```

```

ame.
|
| The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
| existing request. This allows you to change the request for some
| parameters and not others.
|
| .. versionadded:: 1.3
|
| .. note::
|     This method is only relevant if this estimator is used as a
|     sub-estimator of a meta-estimator, e.g. used inside a
|     :class:`pipeline.Pipeline`. Otherwise it has no effect.
|
| Parameters
| -----
| sample_weight : str, True, False, or None,                default=sklearn.utils.metadata_routing.UNCH
ANGED
|     Metadata routing for ``sample_weight`` parameter in ``fit``.
|
| Returns
| -----
| self : object
|     The updated object.
|
| set_inverse_transform_request(self: sklearn.preprocessing._data.StandardScaler, *, copy: Union[bool, NoneType,
str] = '$UNCHANGED$') -> sklearn.preprocessing._data.StandardScaler
|     Request metadata passed to the ``inverse_transform`` method.
|
| Note that this method is only relevant if
| ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
| Please see :ref:`User Guide <metadata_routing>` on how the routing
| mechanism works.
|
| The options for each parameter are:
|
| - ``True``: metadata is requested, and passed to ``inverse_transform`` if provided. The request is ignored
if metadata is not provided.
|
| - ``False``: metadata is not requested and the meta-estimator will not pass it to ``inverse_transform``.
|
| - ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
|
| - ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original n
ame.
|
| The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
| existing request. This allows you to change the request for some
| parameters and not others.
|
| .. versionadded:: 1.3
|
| .. note::
|     This method is only relevant if this estimator is used as a
|     sub-estimator of a meta-estimator, e.g. used inside a
|     :class:`pipeline.Pipeline`. Otherwise it has no effect.
|
| Parameters
| -----
| copy : str, True, False, or None,                default=sklearn.utils.metadata_routing.UNCHANGED
|     Metadata routing for ``copy`` parameter in ``inverse_transform``.
|
| Returns
| -----
| self : object
|     The updated object.
|
| set_partial_fit_request(self: sklearn.preprocessing._data.StandardScaler, *, sample_weight: Union[bool, NoneTy
pe, str] = '$UNCHANGED$') -> sklearn.preprocessing._data.StandardScaler
|     Request metadata passed to the ``partial_fit`` method.
|
| Note that this method is only relevant if
| ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
| Please see :ref:`User Guide <metadata_routing>` on how the routing
| mechanism works.
|
| The options for each parameter are:
|
| - ``True``: metadata is requested, and passed to ``partial_fit`` if provided. The request is ignored if me
tadata is not provided.

```

```

- ``False``: metadata is not requested and the meta-estimator will not pass it to ``partial_fit``.
- ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
existing request. This allows you to change the request for some
parameters and not others.

.. versionadded:: 1.3

.. note::
    This method is only relevant if this estimator is used as a
    sub-estimator of a meta-estimator, e.g. used inside a
    :class:`pipeline.Pipeline`. Otherwise it has no effect.

Parameters
-----
sample_weight : str, True, False, or None,          default=sklearn.utils.metadata_routing.UNCHANGED
    Metadata routing for ``sample_weight`` parameter in ``partial_fit``.

Returns
-----
self : object
    The updated object.

set_transform_request(self: sklearn.preprocessing._data.StandardScaler, *, copy: Union[bool, NoneType, str] =
'UNCHANGED$') -> sklearn.preprocessing._data.StandardScaler
    Request metadata passed to the ``transform`` method.

Note that this method is only relevant if
``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
Please see :ref:`User Guide <metadata_routing>` on how the routing
mechanism works.

The options for each parameter are:

- ``True``: metadata is requested, and passed to ``transform`` if provided. The request is ignored if meta
data is not provided.
- ``False``: metadata is not requested and the meta-estimator will not pass it to ``transform``.
- ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (``sklearn.utils.metadata_routing.UNCHANGED``) retains the
existing request. This allows you to change the request for some
parameters and not others.

.. versionadded:: 1.3

.. note::
    This method is only relevant if this estimator is used as a
    sub-estimator of a meta-estimator, e.g. used inside a
    :class:`pipeline.Pipeline`. Otherwise it has no effect.

Parameters
-----
copy : str, True, False, or None,          default=sklearn.utils.metadata_routing.UNCHANGED
    Metadata routing for ``copy`` parameter in ``transform``.

Returns
-----
self : object
    The updated object.

transform(self, X, copy=None)
    Perform standardization by centering and scaling.

Parameters
-----
X : {array-like, sparse matrix of shape (n_samples, n_features)}
    The data used to scale along the features axis.

```

```

    copy : bool, default=None
        Copy the input X or not.

    Returns
    -----
    X_tr : {ndarray, sparse matrix} of shape (n_samples, n_features)
        Transformed array.

```

Data and other attributes defined here:

```

__annotations__ = {'_parameter_constraints': <class 'dict'>}

```

Methods inherited from sklearn.base.OneToOneFeatureMixin:

```

get_feature_names_out(self, input_features=None)
    Get output feature names for transformation.

    Parameters
    -----
    input_features : array-like of str or None, default=None
        Input features.

    - If `input_features` is `None`, then `feature_names_in_` is
      used as feature names in. If `feature_names_in_` is not defined,
      then the following input feature names are generated:
      `["x0", "x1", ..., "x(n_features_in_ - 1)"]`.
    - If `input_features` is an array-like, then `input_features` must
      match `feature_names_in_` if `feature_names_in_` is defined.

    Returns
    -----
    feature_names_out : ndarray of str objects
        Same as input features.

```

Data descriptors inherited from sklearn.base.OneToOneFeatureMixin:

```

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)

```

Methods inherited from sklearn.base.TransformerMixin:

```

fit_transform(self, X, y=None, **fit_params)
    Fit to data, then transform it.

    Fits transformer to `X` and `y` with optional parameters `fit_params`
    and returns a transformed version of `X`.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Input samples.

    y : array-like of shape (n_samples,) or (n_samples, n_outputs),
        Target values (None for unsupervised transformations).
        default=None

    **fit_params : dict
        Additional fit parameters.

    Returns
    -----
    X_new : ndarray array of shape (n_samples, n_features_new)
        Transformed array.

```

Methods inherited from sklearn.utils._set_output._SetOutputMixin:

```

set_output(self, *, transform=None)
    Set output container.

    See :ref:`sphx_glr_auto_examples_misellaneous_plot_set_output.py`
    for an example on how to use the API.

```


Parameters

`transform` : {"default", "pandas"}, default=None
Configure output of ``transform`` and ``fit_transform``.

- ``default``: Default output format of a transformer
- ``pandas``: DataFrame output
- ``None``: Transform configuration is unchanged

Returns

`self` : estimator instance
Estimator instance.

Class methods inherited from `sklearn.utils._set_output._SetOutputMixin`:

`__init_subclass__`(auto_wrap_output_keys=('transform',), **kwargs) from `builtins.type`
This method is called when a class is subclassed.

The default implementation does nothing. It may be overridden to extend subclasses.

Methods inherited from `sklearn.base.BaseEstimator`:

`__getstate__`(self)
Helper for pickle.

`__repr__`(self, N_CHAR_MAX=700)
Return `repr(self)`.

`__setstate__`(self, state)

`__sklearn_clone__`(self)

`get_params`(self, deep=True)
Get parameters for this estimator.

Parameters

`deep` : bool, default=True
If True, will return the parameters for this estimator and contained subobjects that are estimators.

Returns

`params` : dict
Parameter names mapped to their values.

`set_params`(self, **params)
Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as `:class:`~sklearn.pipeline.Pipeline``). The latter have parameters of the form ``<component>__<parameter>`` so that it's possible to update each component of a nested object.

Parameters

`**params` : dict
Estimator parameters.

Returns

`self` : estimator instance
Estimator instance.

Methods inherited from `sklearn.utils._metadata_requests._MetadataRequester`:

`get_metadata_routing`(self)
Get metadata routing of this object.

Please check `:ref:`User Guide <metadata_routing>`` on how the routing mechanism works.

Returns

```
|         routing : MetadataRequest
|         A :class:`~utils.metadata_routing.MetadataRequest` encapsulating
|         routing information.
```

```
In [15]: # toy data
train = {'salary':[50_000,75_000,40_000,1_000_000,30_000,250_000,35_000,45_000]}
test = {'salary':[25_000,55_000,1_500_000,60_000]}

Xtoy_train = pd.DataFrame(train)
Xtoy_test = pd.DataFrame(test)

scaler = StandardScaler()
print(scaler.fit_transform(Xtoy_train))
print(scaler.transform(Xtoy_test))

[[-0.44873188]
 [-0.36895732]
 [-0.4806417 ]
 [ 2.58270127]
 [-0.51255153]
 [ 0.18946457]
 [-0.49659661]
 [-0.46468679]]
[[-0.52850644]
 [-0.43277697]
 [ 4.1781924 ]
 [-0.41682206]]
```

```
In [16]: # adult data

std_ftrs = ['capital-gain', 'capital-loss']
scaler = StandardScaler()
print(scaler.fit_transform(X_train[std_ftrs]))
print(scaler.transform(X_val[std_ftrs]))
print(scaler.transform(X_test[std_ftrs]))

[[-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 ...
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]]
[[-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 ...
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]]
[[-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 ...
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]
 [-0.14633293 -0.22318878]]
```

Quiz 2


Which of these features could be safely preprocessed by the minmax scaler?

- number of minutes spent on the website in a day
- number of days a year spent abroad in a year
- USD donated to charity

How and when to do preprocessing in the ML pipeline?

- **APPLY TRANSFORMER.FIT ONLY ON YOUR TRAINING DATA!** Then transform the validation and test sets.
- One of the most common mistake practitioners make is leaking statistics!
 - `fit_transform` is applied to the whole dataset, then the data is split into train/validation/test
 - this is wrong because the test set statistics impacts how the training and validation sets are transformed
 - but the test set must be separated by train and val, and val must be separated by train

- or fit_transform is applied to the train, then fit_transform is applied to the validation set, and fit_transform is applied to the test set
 - this is wrong because the relative position of the points change



Scikit-learn's pipelines

- The steps in the ML pipeline can be chained together into a scikit-learn pipeline which consists of transformers and one final estimator which is usually your classifier or regression model.
- It neatly combines the preprocessing steps and it helps to avoid leaking statistics.

https://scikit-learn.org/stable/auto_examples/compose/plot_column_transformer_mixed_types.html

```
In [17]: import pandas as pd
import numpy as np

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split

#np.random.seed(0)

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k or less than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X, y, train_size = 0.6, random_state=random_state)

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other, y_other, train_size = 0.5, random_state=random_state)

In [18]: # collect which encoder to use on each feature
# needs to be done manually
ordinal_ftrs = ['education']
ordinal_cats = [['Preschool', '1st-4th', '5th-6th', '7th-8th', '9th', '10th', '11th', '12th', 'HS-grad', \
                'Some-college', 'Assoc-voc', 'Assoc-acdm', 'Bachelors', 'Masters', 'Prof-school', 'Doctorate']]
onehot_ftrs = ['workclass', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'native-country']
minmax_ftrs = ['age', 'hours-per-week']
std_ftrs = ['capital-gain', 'capital-loss']

# collect all the encoders
preprocessor = ColumnTransformer(
    transformers=[
        ('ord', OrdinalEncoder(categories = ordinal_cats), ordinal_ftrs),
        ('onehot', OneHotEncoder(sparse=False, handle_unknown='ignore'), onehot_ftrs),
        ('minmax', MinMaxScaler(), minmax_ftrs),
        ('std', StandardScaler(), std_ftrs)])

clf = Pipeline(steps=[('preprocessor', preprocessor)]) # for now we only preprocess
                                                    # later on we will add other steps here

X_train_prep = clf.fit_transform(X_train)
X_val_prep = clf.transform(X_val)
```

```
X_test_prep = clf.transform(X_test)
```

```
print(X_train.shape)
print(X_train_prep.shape)
print(X_train_prep)
```

```
(19536, 14)
(19536, 91)
[[10.          0.          0.          ...  0.39795918 -0.14633293
  -0.22318878]
 [ 9.          0.          0.          ...  0.39795918 -0.14633293
  -0.22318878]
 [ 8.          0.          0.          ...  0.5          -0.14633293
  -0.22318878]
 ...
 [ 6.          0.          0.          ...  0.19387755 -0.14633293
  -0.22318878]
 [ 8.          0.          0.          ...  0.84693878 -0.14633293
  -0.22318878]
 [12.         0.          0.          ...  0.60204082 -0.14633293
  -0.22318878]]
```

```
/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:972: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
  warnings.warn(
```

Mudcard

In []: